



**SOFTWARE
AND SYSTEMS
ENGINEERING**
research group

Model-based testing under uncertainty

L1: Theoretical aspects & practical applications

Matteo Camilli

matteo.camilli@unibz.it

<https://matteocamilli.github.io>

Formal Methods at Work @
Gran Sasso Science Institute (GSSI)
A.Y. 2019/20



Outline

- Markov Decision Process
 - Structure
 - Rewards
 - Policy, best policy, value iteration
- Model-based testing (MBT)
 - Offline vs online approaches
 - Conformance relation
 - Probabilistic alternating simulation and refinement
- Online MBT under uncertainty
 - Problem statement
 - Uncertain model parameters
 - Bayesian inference
 - Framework and test case generation strategies

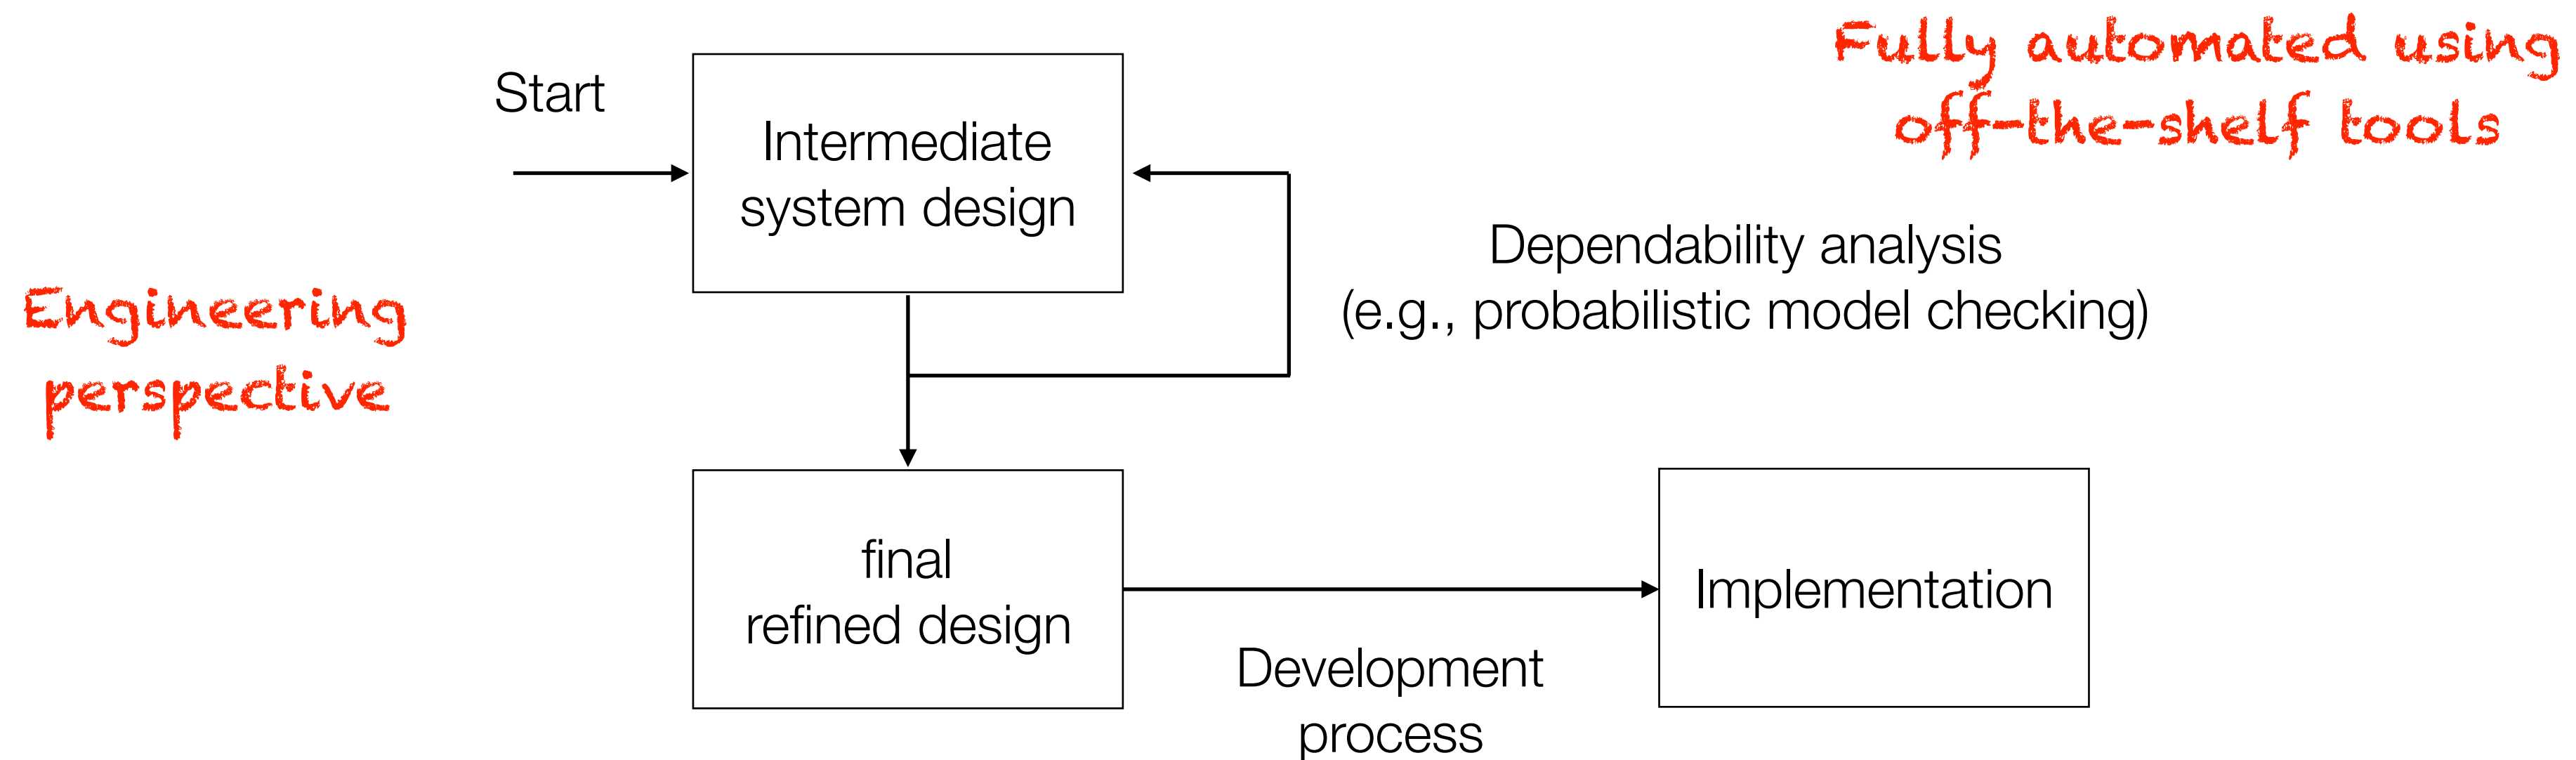
Markov decision processes

- Structure
- Examples
- Rewards
- Policy, best policy
- Value iteration algorithms
- Probabilistic model checking with PCTL

Markov Models

- **Basic notions**

- Formal framework for performance and dependability (reliability, availability, safety) analysis
- Dependability modeling (upfront) at design time improves the quality of the system eventually produced
- Assumption — the modeled system meets the Markov property^{1,2} (memoryless)



1. The the probability of moving to the next state only depends on the current state, not on the history that lead to that state.

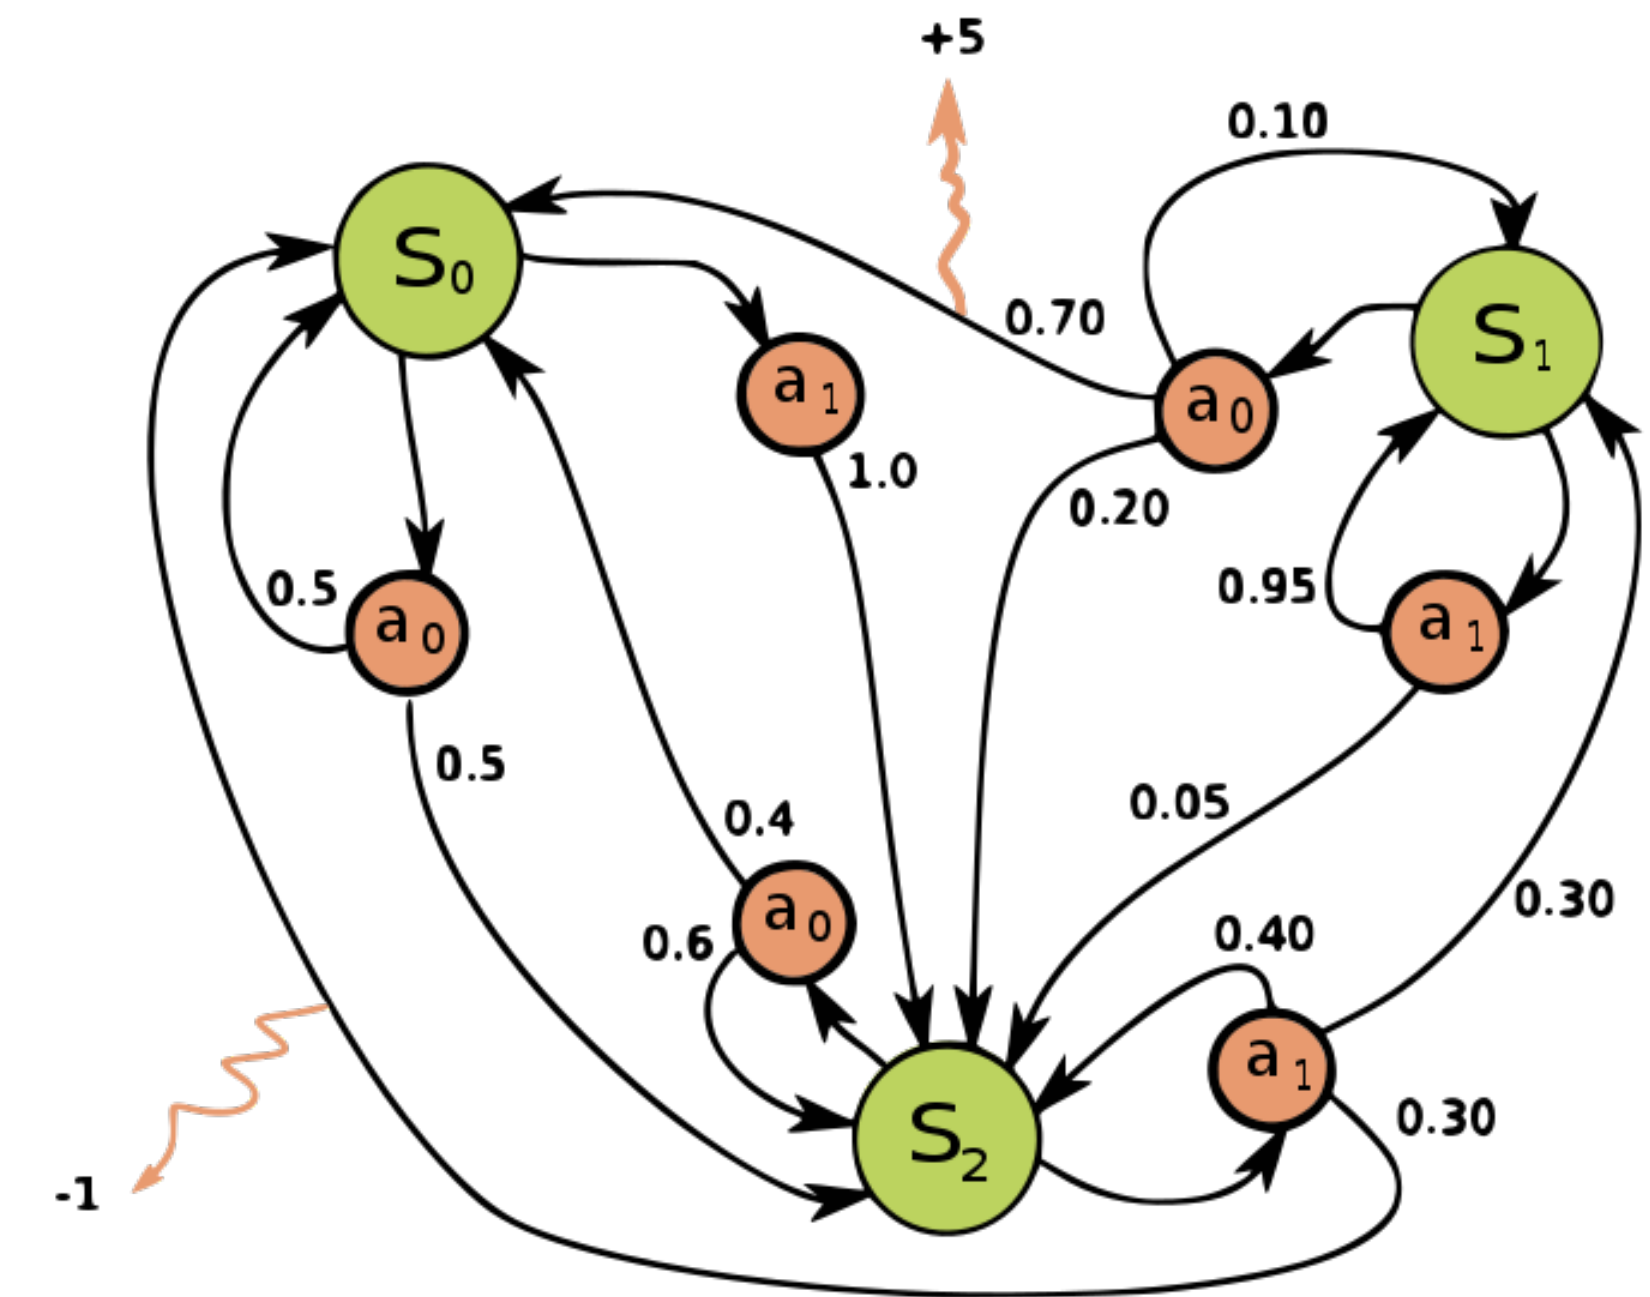
2. R. C. Cheung, A user-oriented software reliability model, IEEE TSE, no. 2, pp. 118– 125, 1980

Markov Decision Process

- Mathematical framework for modeling systems whose behavior is partially
 - Nondeterministic — actions (external stimuli) under the control of a decision maker
 - Stochastic — random outcome out of an executed action

- **Formal structure**

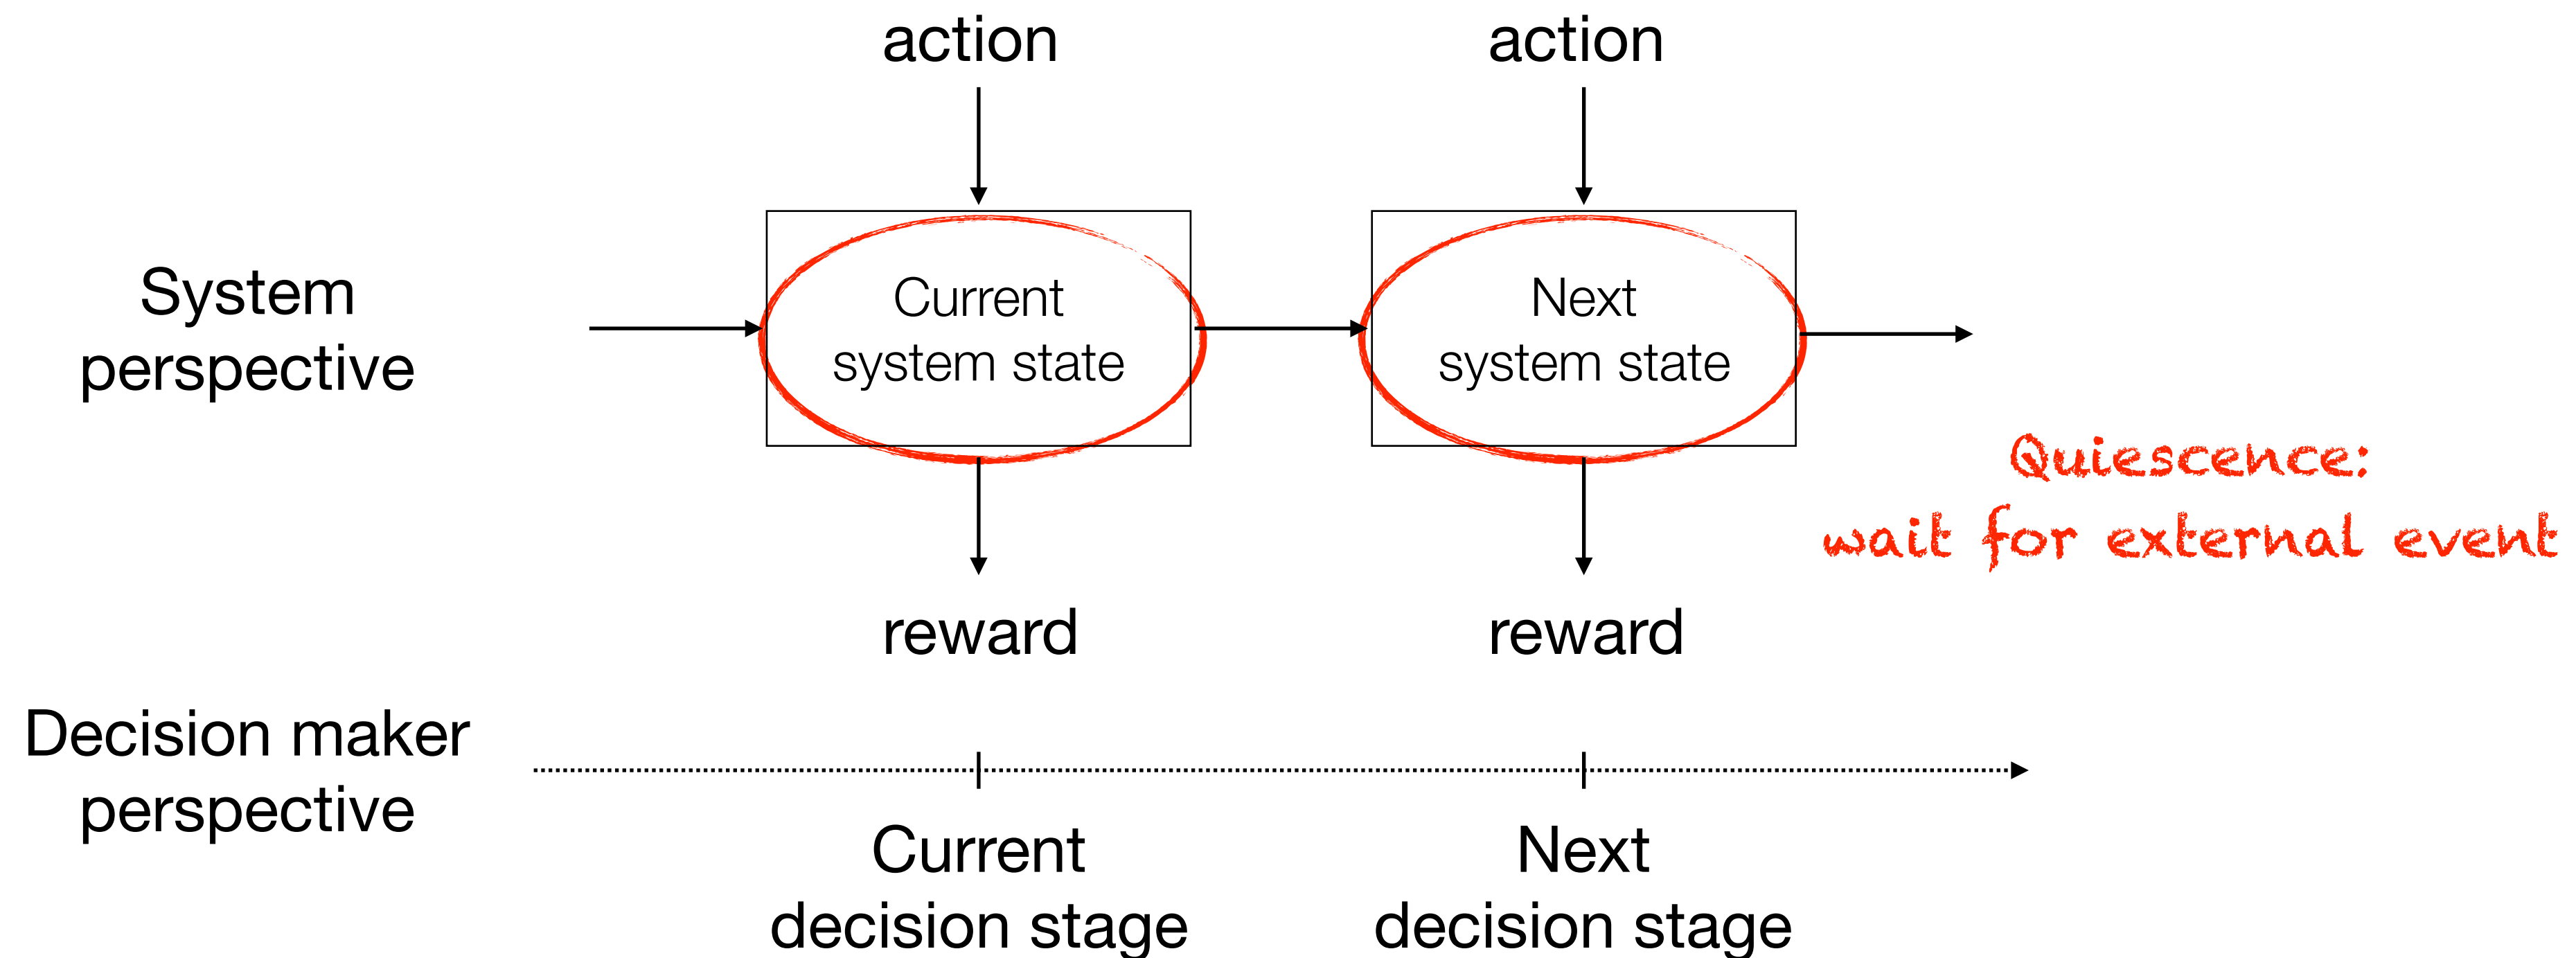
- S : set of states
- s_0 : initial state
- A : set of actions
- $P: S \times A \times S \rightarrow [0,1], P(s, a, s') = P(s_{t+1} = s' \mid s_t = s, a_t = a)$
- $R: S \times A \times S \rightarrow \mathbb{R}, R(s, a, s')$ reward for $(s_{t+1} = s', s_t = s, a_t = a)$



MDP behavior

- How the model operates

- The system must be in one of the states (finite countable set) at a time
- The system makes a transition $s \rightarrow s'$ when one of the available actions is selected



MDP behavior

- **States**

- System configurations or operational status of components
- Instances of the system where
 - Components are operational or failed (e.g., enumeration of working/failed components)
 - Experienced specific sequences of events (e.g., events observed so far)
 - Undergoing recover/repair
 - Operating in a fully-functioning mode, degraded mode, faulty, etc.

- **Actions**

- Possible inputs or external events

- **Transitions**

- Define whether is possible to go from one state to another
- Transition probability \rightarrow governs the likelihood of observing the transition

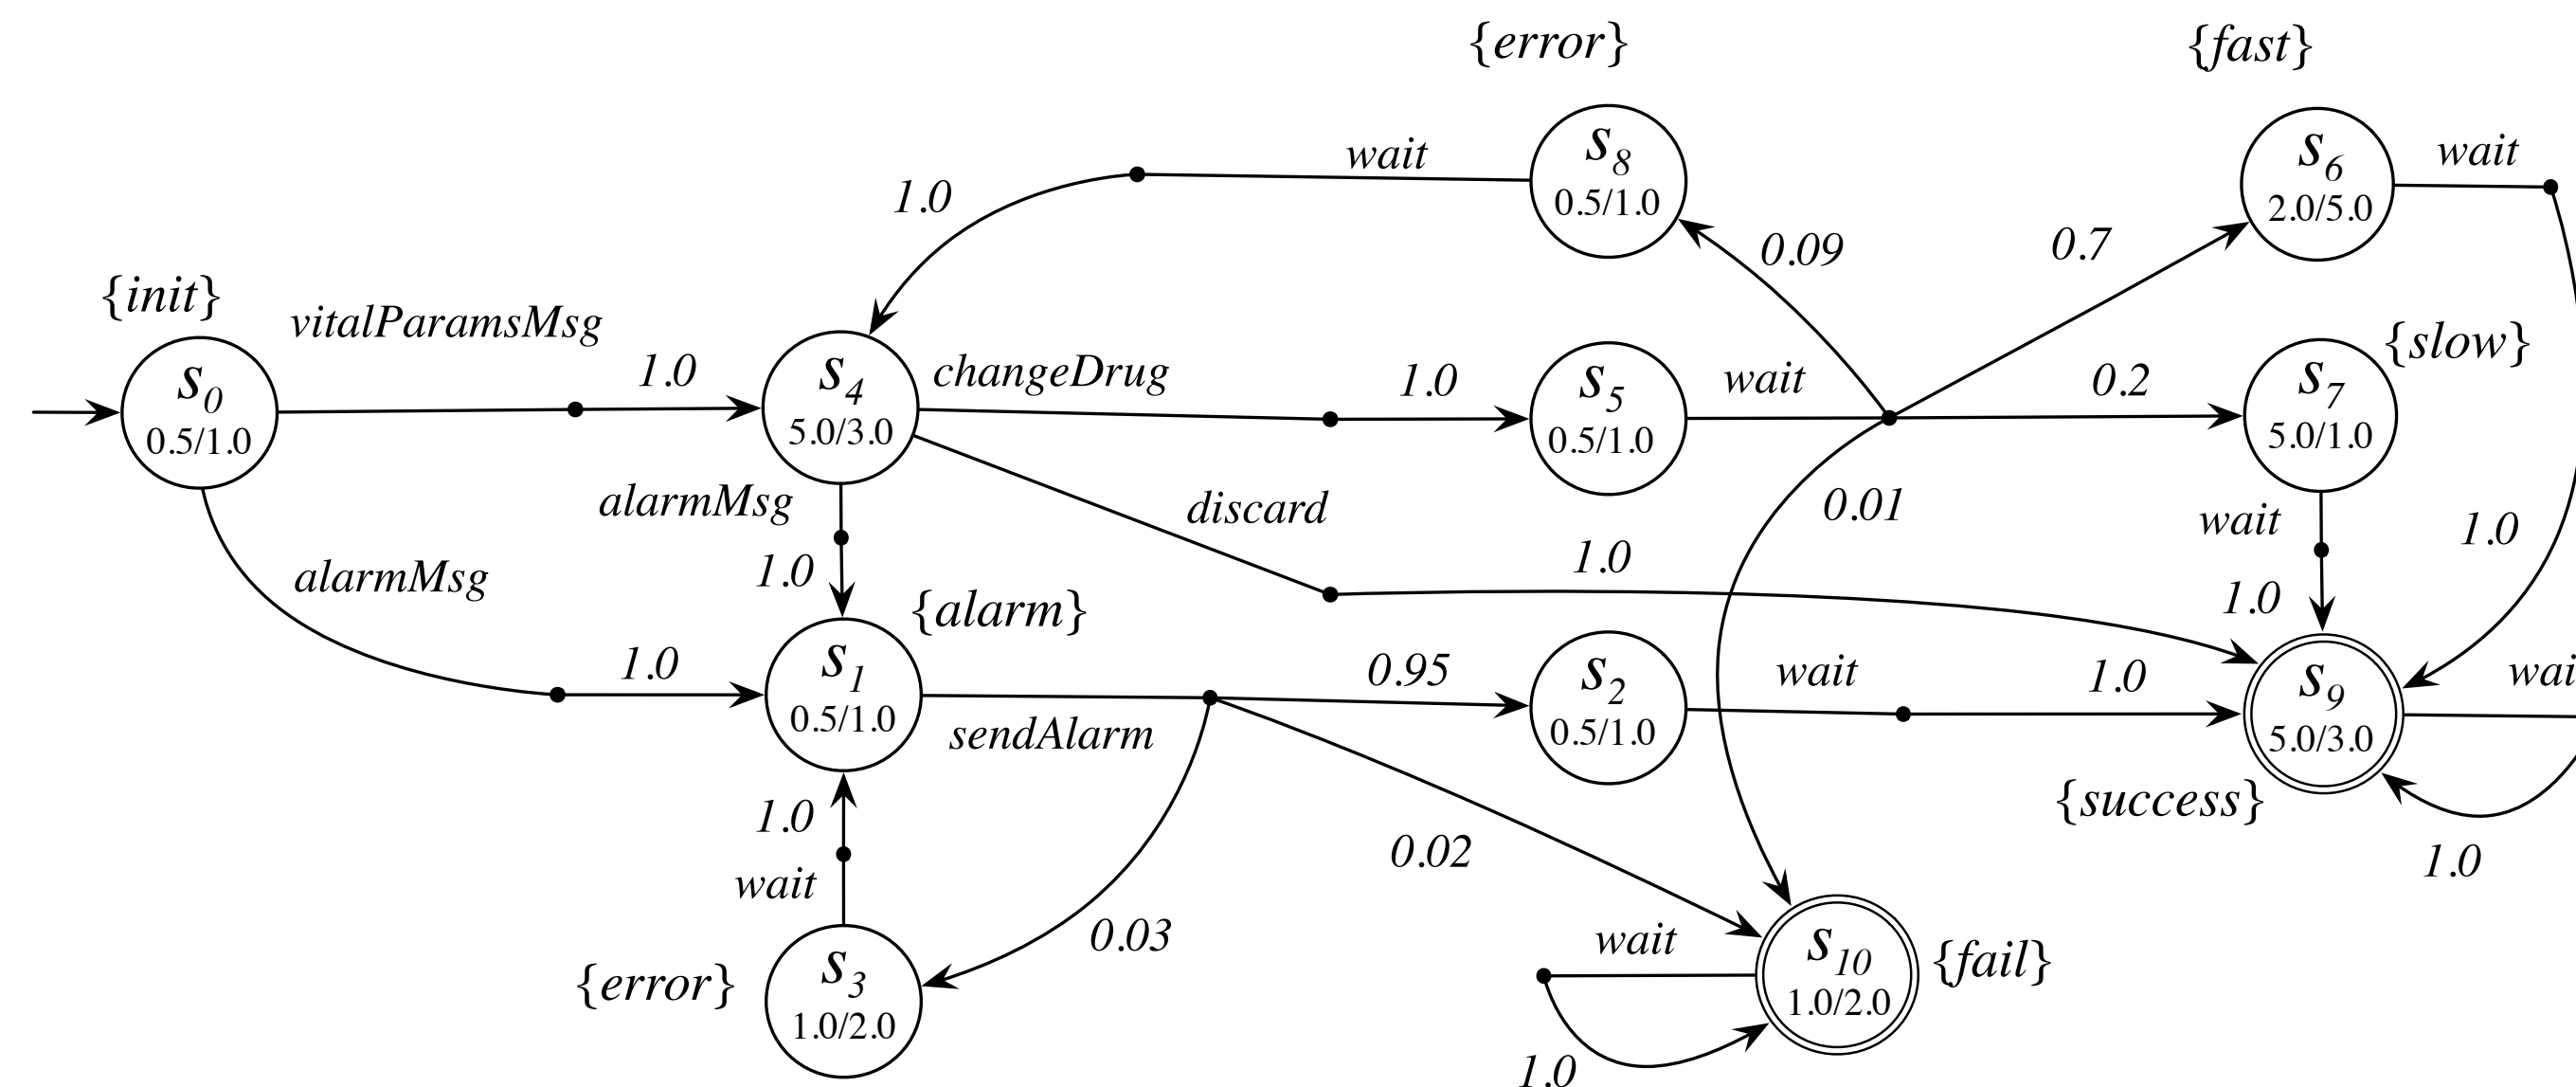
MDP examples

- Possible scenarios
 - Service-based systems
 - Web applications
 - Mobile applications
 - Cyber physical systems (CPSs)
 - etc.
- Selected examples
 - Tele assistant system (TAS) — example of service-based system
 - SafeHome — example of CPS

MDP examples — TAS

- **TAS**^{1,2}

- SBS providing health support to chronic condition patients at their homes
- wearable devices (track vital parameters) + remote services (healthcare, pharmacy and emergency units)



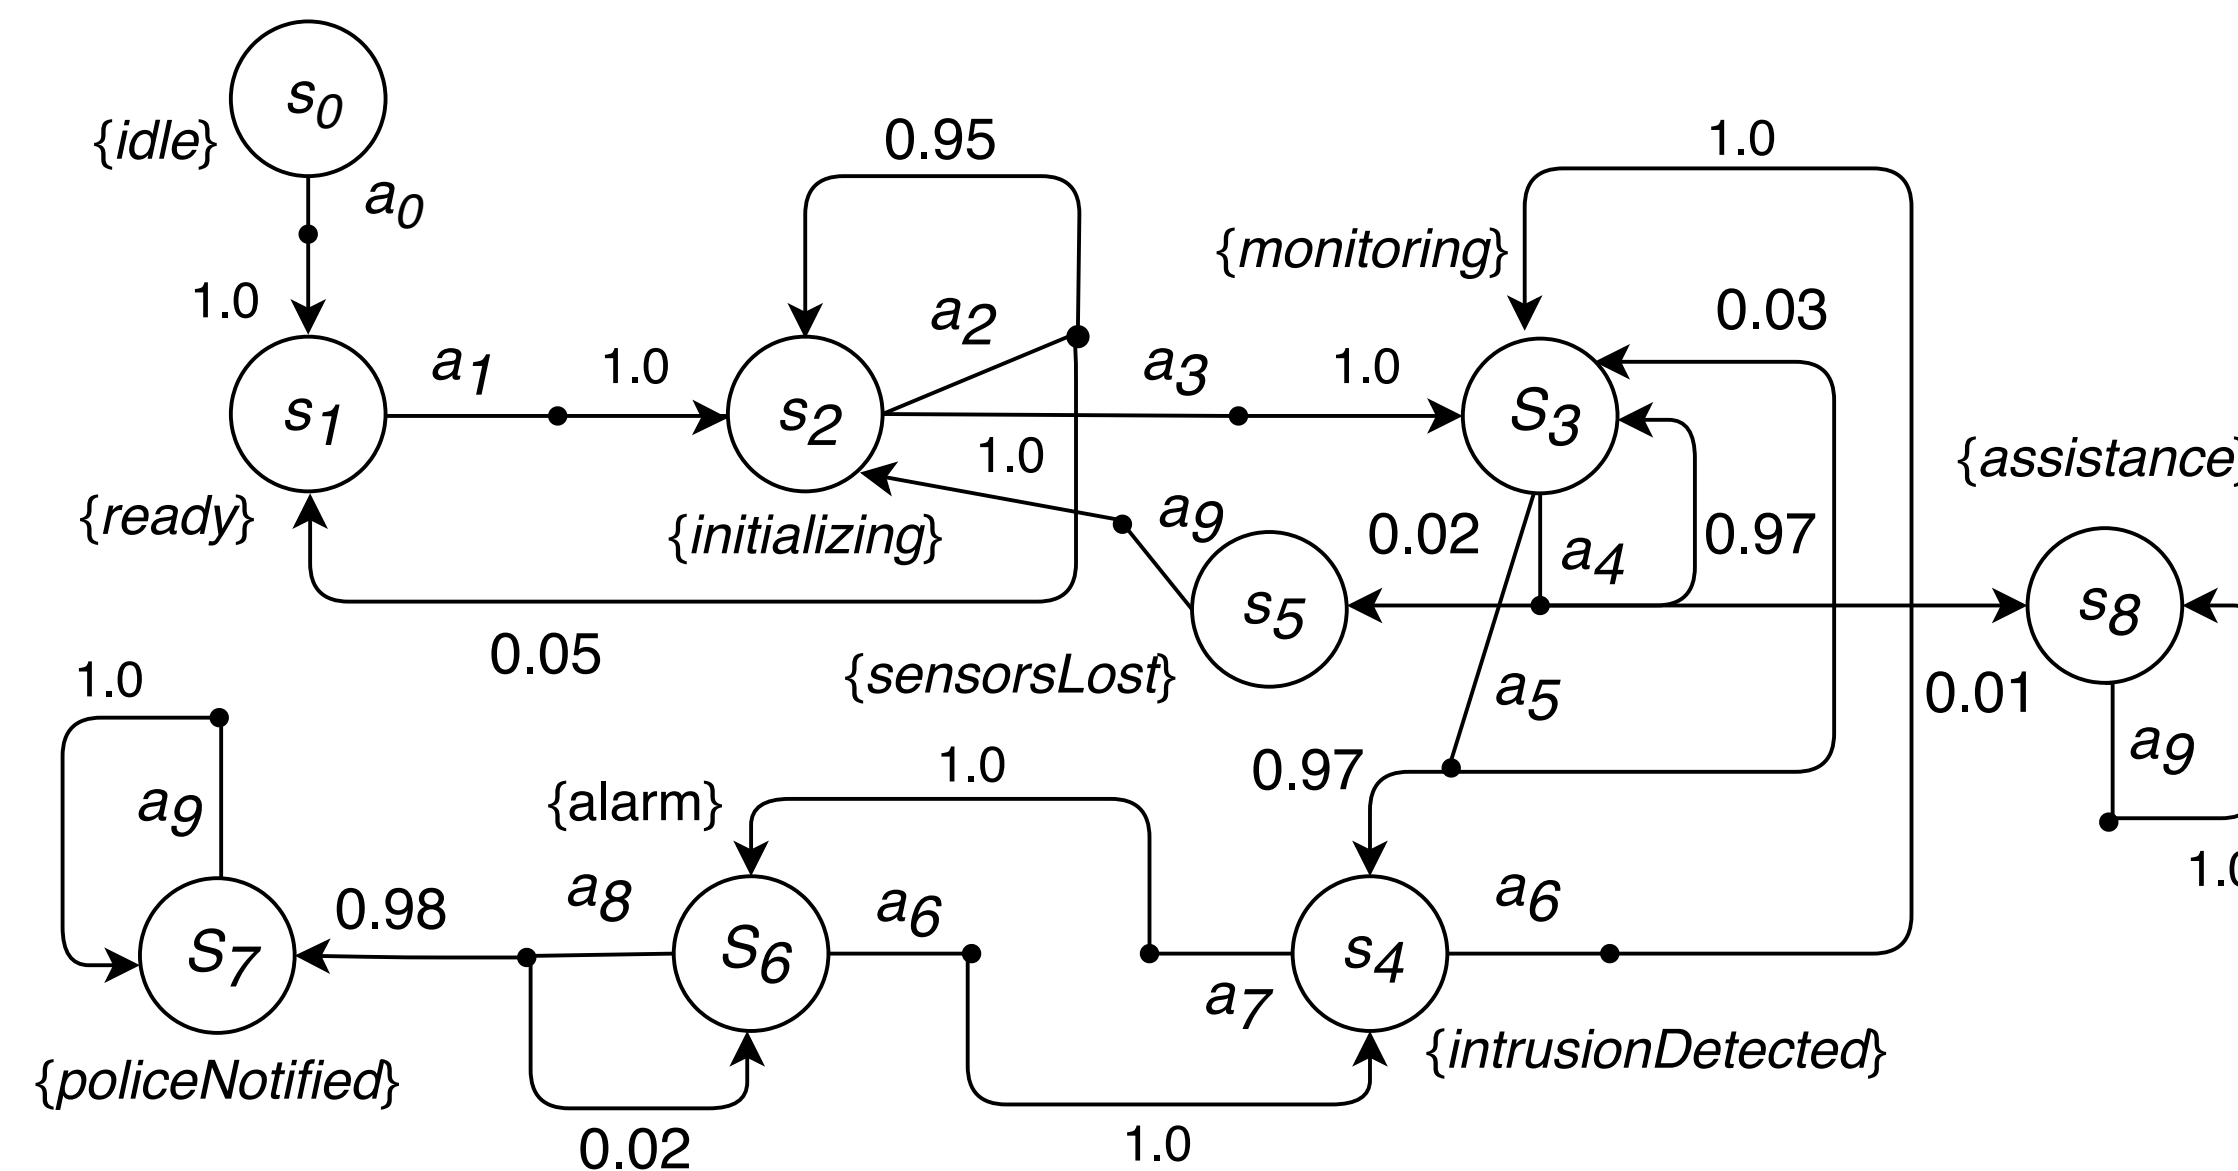
1. D. Weyns et al., Tele Assistance: A Self-Adaptive Service-Based System Exemplar, SEAMS 2015, IEEE

2. M. Camilli et al., Online Model-Based Testing under Uncertainty, ISSRE 2018, IEEE

MDP examples — SafeHome

- **SafeHome**^{1,2}

- Configure and control alarms along with related sensors that implement security and safety features
- Here is the description of a part in charge of monitoring and intrusion detection
 - Modeled phases: sensors initialization, monitoring, detection, notification



Actions

- a0: activate
- a1: startInit
- a2: initSensors
- a3: startMonitoring
- a4: sensorsCheck
- a5: intrusionOccurred
- a6: cancel
- a7: turnAlarmOn
- a8: notify
- a9: wait

1. Roger S Pressman. 2005. Software engineering: a practitioner's approach. Palgrave Macmillan

2. Man Zhang, et al., Uncertainty-wise test case generation and minimization for Cyber-Physical Systems. 2019, JSS 153

MDP with rewards

- An MDP model can be augmented with multiple reward structures
- **Reward structure**
 - $R: S \times A \times S \rightarrow \mathbb{R}$, $R(s, a, s')$ reward for $(s_{t+1} = s', s_t = s, a_t = a)$
 - Describe nonfunctional aspects (e.g., energy consumption, computational cost, response time, ...)



Mental note

This is the usual interpretation of a reward structure.
We'll see how to leverage this notion in a “unconventional” way to drive testing.

MDP policy

- The notion of policy π refers to the way a Decision Maker (DM) solves nondeterminism of a MDP
- **Deterministic policy**^{1,2}
 - $\pi: S \rightarrow A$, prescribes the action to take given a state
 - DM objective: choose π which maximizes the expected cumulated reward over an infinite horizon
 - This is called best policy π^*
- **Computation of the best deterministic policy**
 - Given $R(s, a)$, i.e., the one-step expected reward
 - We can compute the value function $V(s)$ for each state

Bellman's equation

$$R(s, a) = \sum_{s' \in S} p_{s,a,s'} r_{s,a,s'} \quad V(s) = \max_{a \in A} \{R(s, a) + \gamma \sum_{s' \in S} p_{s,a,s'} V(s')\}$$

Best policy

$$\pi^*(s) = \arg \max_{a \in A} \{R(s, a) + \gamma \sum_{s' \in S} p_{s,a,s'} V(s')\}$$

1. Given a deterministic policy, the MDP reduces to a Discrete Time Markov Chain (DTMC).

2. Martin L. Puterman. 1994. Markov Decision Processes: Discrete Stochastic Dynamic Programming (1st. ed.). John Wiley & Sons, Inc., USA

MDP policy — Value iteration

```

1: Procedure ValueIteration( $S, A, P, R, \theta$ )
2:   Inputs
3:      $S$  set of all states
4:      $A$  set of all actions
5:      $P$  transition function  $P(s, a, s')$ 
6:      $R$  reward function  $R(s, a, s')$ 
7:      $\theta$  a threshold,  $\theta > 0$ 
8:   Output
9:      $\pi[S]$  optimal policy
10:     $V[S]$  value function
11:   Local
12:     real array  $V_k[S]$  is a sequence of value functions
13:     action array  $\pi[S]$ 
14:   assign  $V_0[S]$  arbitrarily
15:    $k \leftarrow 0$ 
16:   repeat
17:      $k \leftarrow k + 1$ 
18:     for each state  $s$  do
19:        $V_k[s] = \max_a \sum_{s'} P(s, a, s') (R(s, a, s') + \gamma V_{k-1}[s'])$ 
20:        $\pi[S] = a$ 
21:   until  $\forall s |V_k[s] - V_{k-1}[s]| < \theta$ 
22:   return  $\pi, V_k$ 

```

- The ValueIteration procedure uses **dynamic programming**
 - Memoization + recursion (or iteration)
- This procedure converges no matter what is the initial value function V_0

(current problem) (subproblem)

Update V_k based on V_{k-1}

Model-based testing of probabilistic systems

- Offline vs online approaches
- Conformance relation
- Probabilistic alternating simulation and refinement

Model-based testing

- **Basic idea**

- A formal model of the required behavior of the System Under Test (SUT) is used as baseline of
 - test case generation
 - Construction of the oracle
- Test suites are automatically extracted from models and then executed

- **Formal verification vs Model-based testing**

- Formal verification — prove that the model (i.e., formal specification) satisfies requirements
- MBT — show that the SUT behaves as defined in the (verified) model
 - Limitation: testing is not complete (i.e., can only show the presence of errors, not their absence)

MBT — terminology

- **Implementation or System Under Test**

- Piece of hardware/software, a software system, an embedded system, a CPS, etc.
- The SUT is viewed as a black-box (secret internal structure)
- The tester controls and observes the SUT via its interfaces (e.g., APIs)

- **Specification**

- Describes what the SUT should do using a formal notation (or language)
- *SPEC* — set of all valid models in a formal notation
 - A specification is $M \in SPEC$

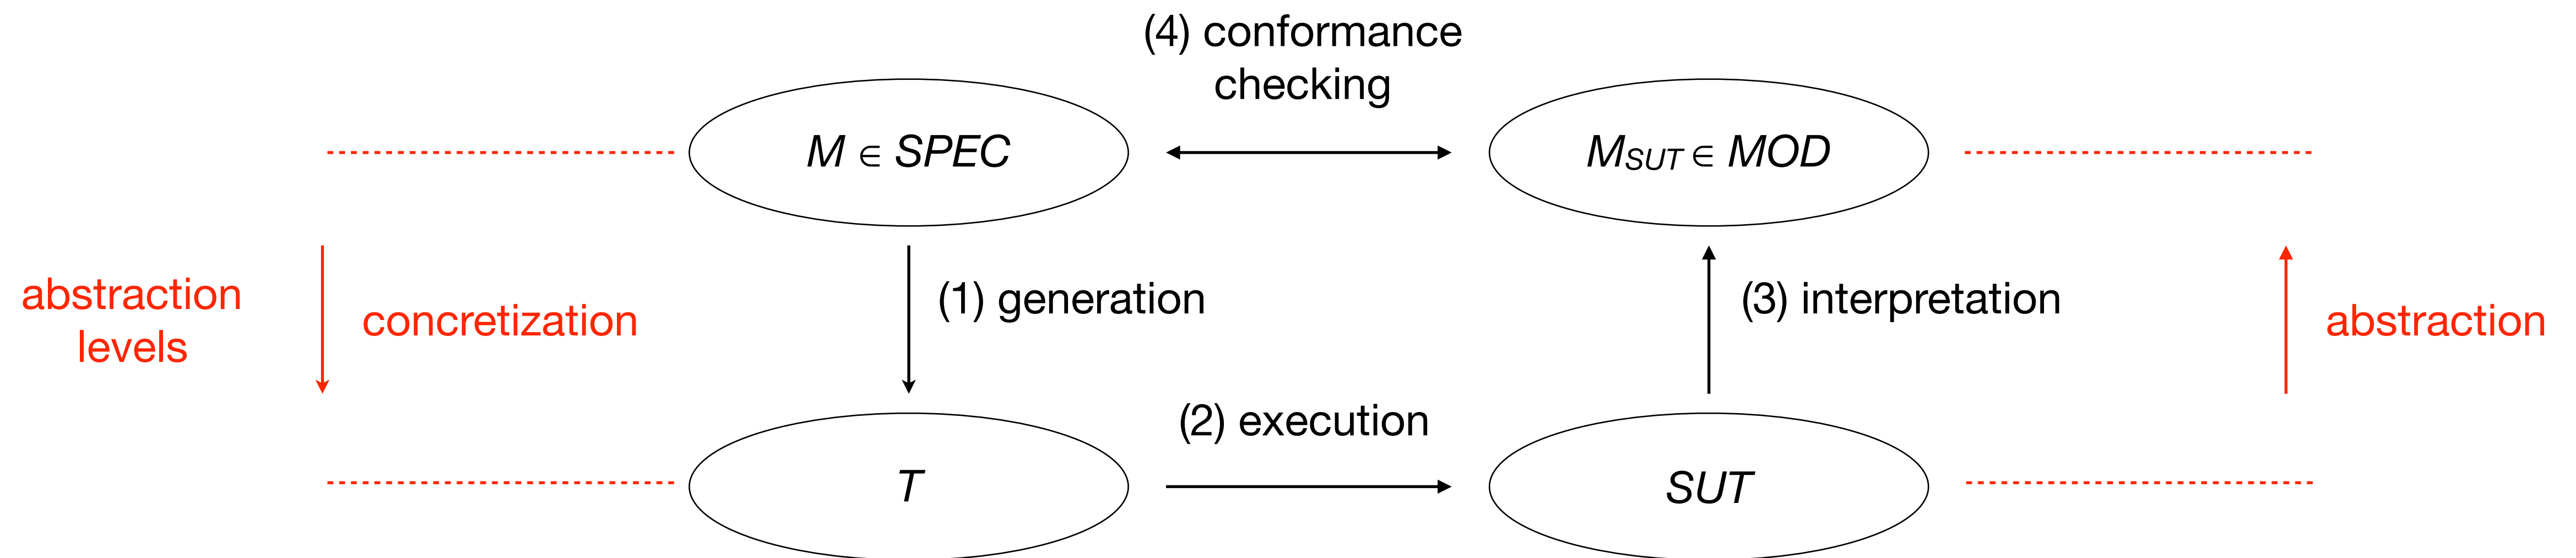
- **Conformance**

- Formalizes the notion of correct behavior of a SUT w.r.t. $M \in SPEC$
- **Problem**: we'd like to define a relation between elements of different domains
 - M (formal object) \longleftrightarrow SUT (not a formal object)

MBT — terminology (2)

- Conformance Problem: M (formal object) \longleftrightarrow SUT (not a formal object)
 - Trick — test assumption
 - The SUT behavior can be interpreted using the same level of abstraction of M
 - The SUT behavior is a model $M_{SUT} \in MOD \subseteq SPEC$
 - MOD — universe of implementation models
 - M_{SUT} not a-priori known
 - Conformance (under the test assumption)
 - Can be expressed as a formal relation between MOD and $SPEC$ elements
 - $conf \subseteq MOD \times SPEC$
 - M_{SUT} is correct w.r.t. M if $M_{SUT} conf S$
 - Conformance checking
 - Assess by testing whether $M_{SUT} conf S$
 - Create T (test suite) s.t. $M_{SUT} conf M \Rightarrow M_{SUT} \text{ passes } T$ (sound but not complete)

MBT process




- Offline vs online ¹
 - Offline — steps 1-4 are separated
 - Online (or on-the-fly) — steps 1-4 are merged into a one-iteration step
 - Test cases are created dynamically and take advantage of the knowledge gained by exploring S
 - Iterative approach —> 2-players game: controller + observer

1. Utting, Mark, and Bruno Legeard. Practical model-based testing: a tools approach. Elsevier, 2010

Online MBT with MDPs

- **Binding** (concretization)

- Defines a mapping between the MDP model spec and the SUT behavior



- **Formal definition**

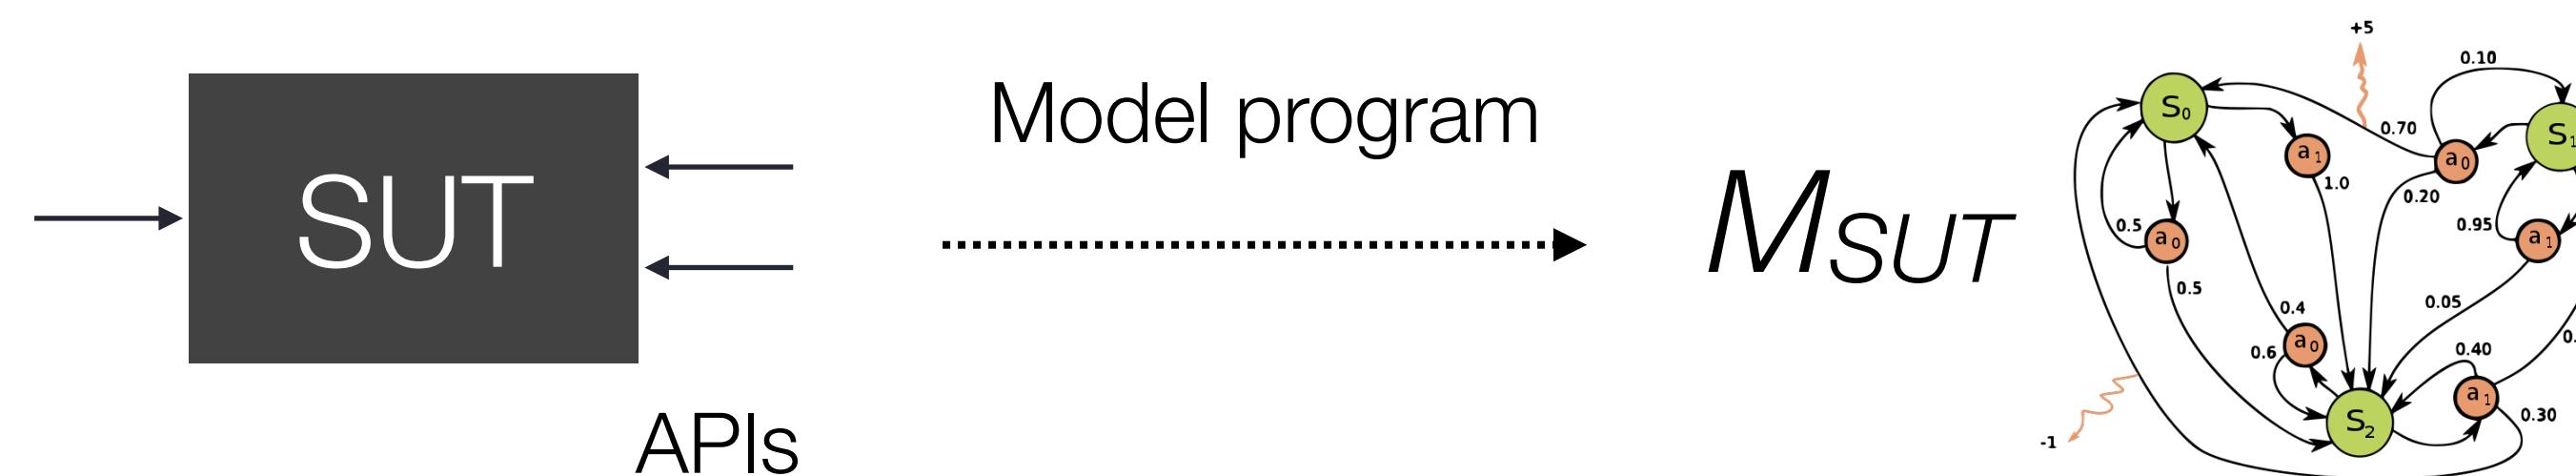
- Given a MDP $M = (S, s_0, A, P)$ and a the SUT with a set of exported services H , a binding is a tuple of partial functions (h, i, post) having domain s.t.
 - $h(s, a), a \in A(s)$ identifies a service $\in H$
 - $i(s, a), a \in A(s)$ identifies a vector v_{in} for the service $h(s, a)$
 - $\text{post}(s, a, s'), a \in A(s)$ maps to a post-condition that must hold for v_{out} resulting from the execution of the service $h(s, a)$ on input v_{in}

Controllable
SUT components

Observable
SUT
components

Online MBT with MDPs (2)

- **Model program** (abstraction)
 - Defines the abstract interpretation of the SUT behavior in terms of MDP model



- **Formal definition**
 - Given a MDP $M = (S, s_0, A, P)$ and a binding (h, i, post) the model program $M_{SUT} = (S', s_0', A', P')$ is a MDP model s.t.
 - $S' \subseteq S, A' \supseteq A$ ← ALL observable SUT states exist in M
 - $s_0 = s_0'$ ← ALL controllable actions in M are feasible in SUT
 - $P'(s, a, s') > 0$ iff $\text{post}(s, a, s')$ holds for v_{out} after $h(s, a)(v_{\text{in}})$ execution

← M_{SUT} transitions are defined in terms of SUT behavior


Online MBT with MDPs (3)

- **Conformance checking**

- Needs the definition of conformance relation = probabilistic alternating simulation + refinement

- **Probabilistic alternating simulation**

- between M and M_{SUT} is a binary relation $\sigma \subseteq S \times S'$, s.t. for all $(s, s') \in \sigma$
 - $A(s) \subseteq A'(s')$
 - For each $t \in S : P(s, a, t) > 0$, there exists $t' \in S' : P'(s', a, t') > 0$ and $(t, t') \in \sigma$



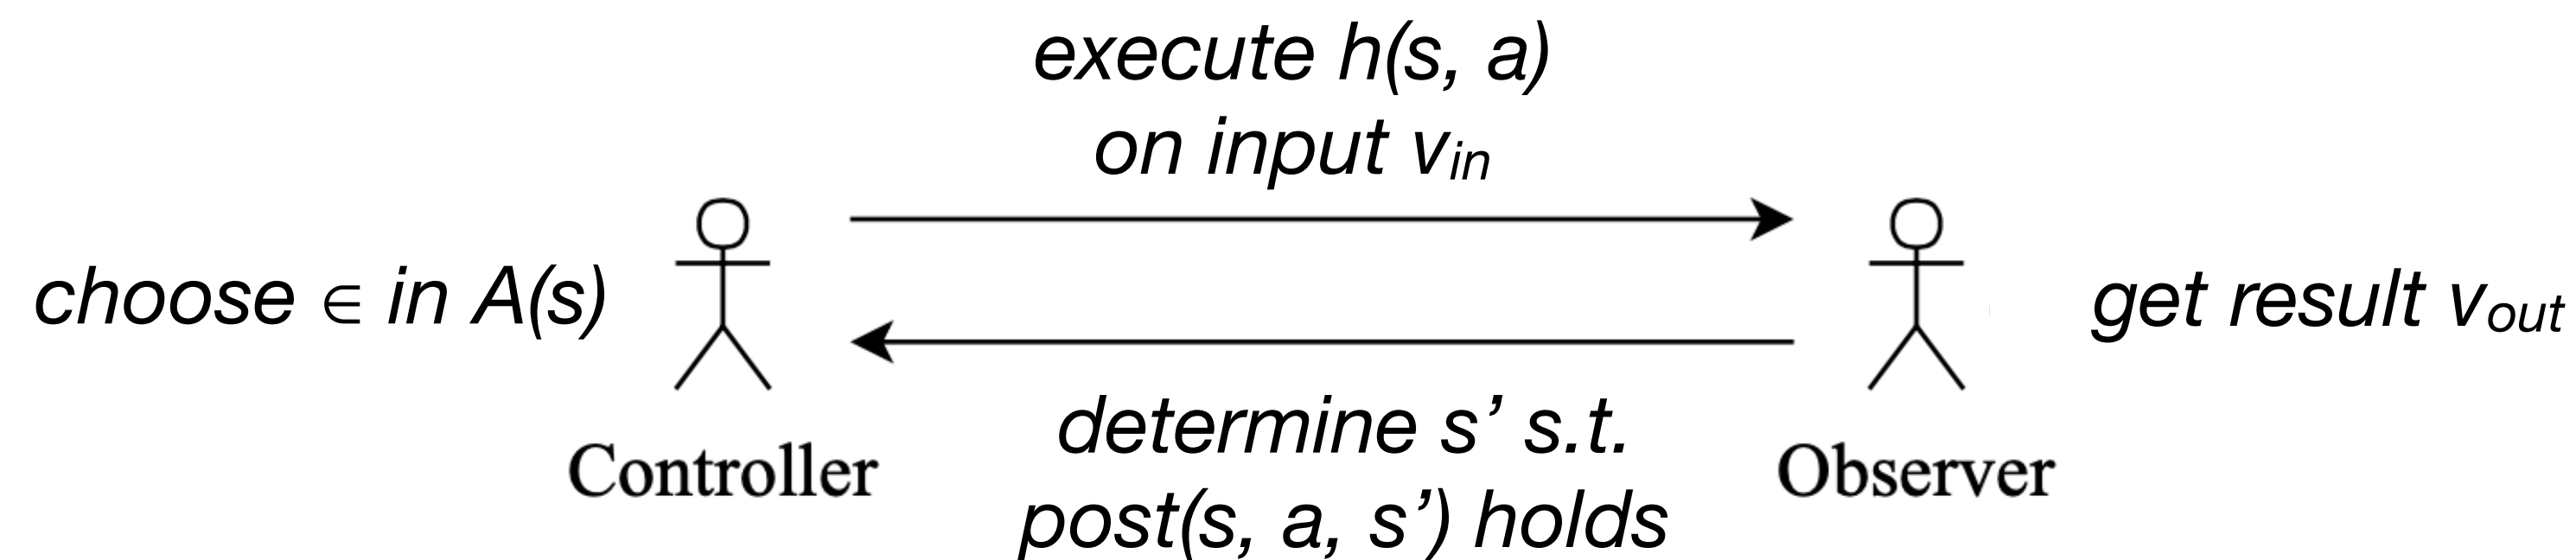
Online MBT with MDPs (4)

• Refinement

- M_{SUT} refines M iff there exists a probabilistic alternating simulation σ s.t. $(s_0, s_0') \in \sigma$

• Conformance game

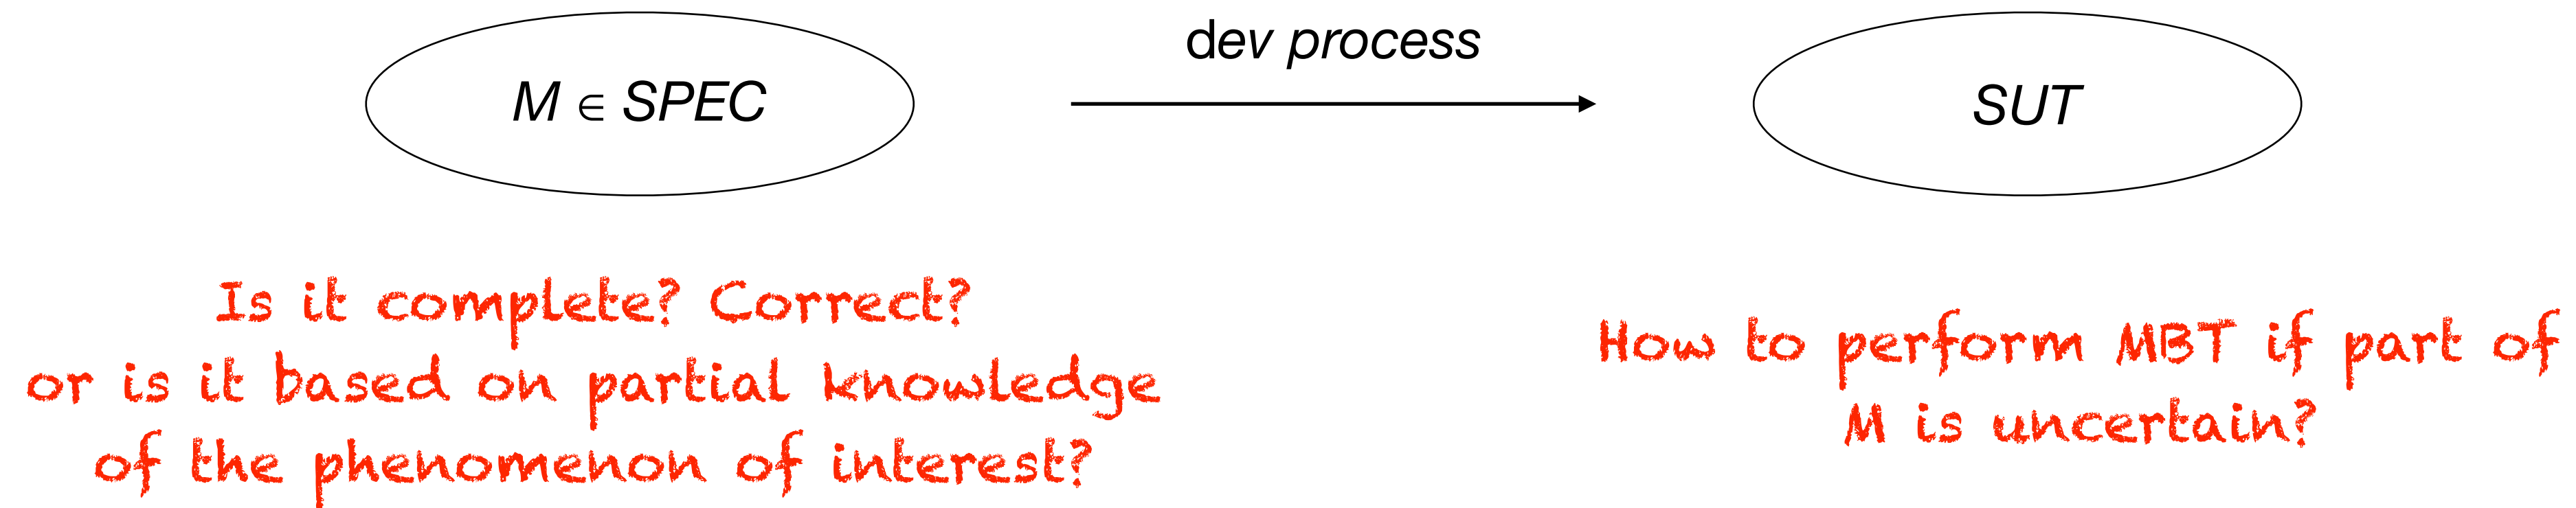
- The notion of refinement is verified in practice by means of a conformance game between
 - Controller \rightarrow chooses actions based on a given test case generation strategy
 - Observer \rightarrow verifies the result out of a test execution



Online MBT under Uncertainty

- Problem statement
- Uncertain model parameters
- Bayesian inference
- Framework and test case generation strategies

Problem statement



- **Problem**

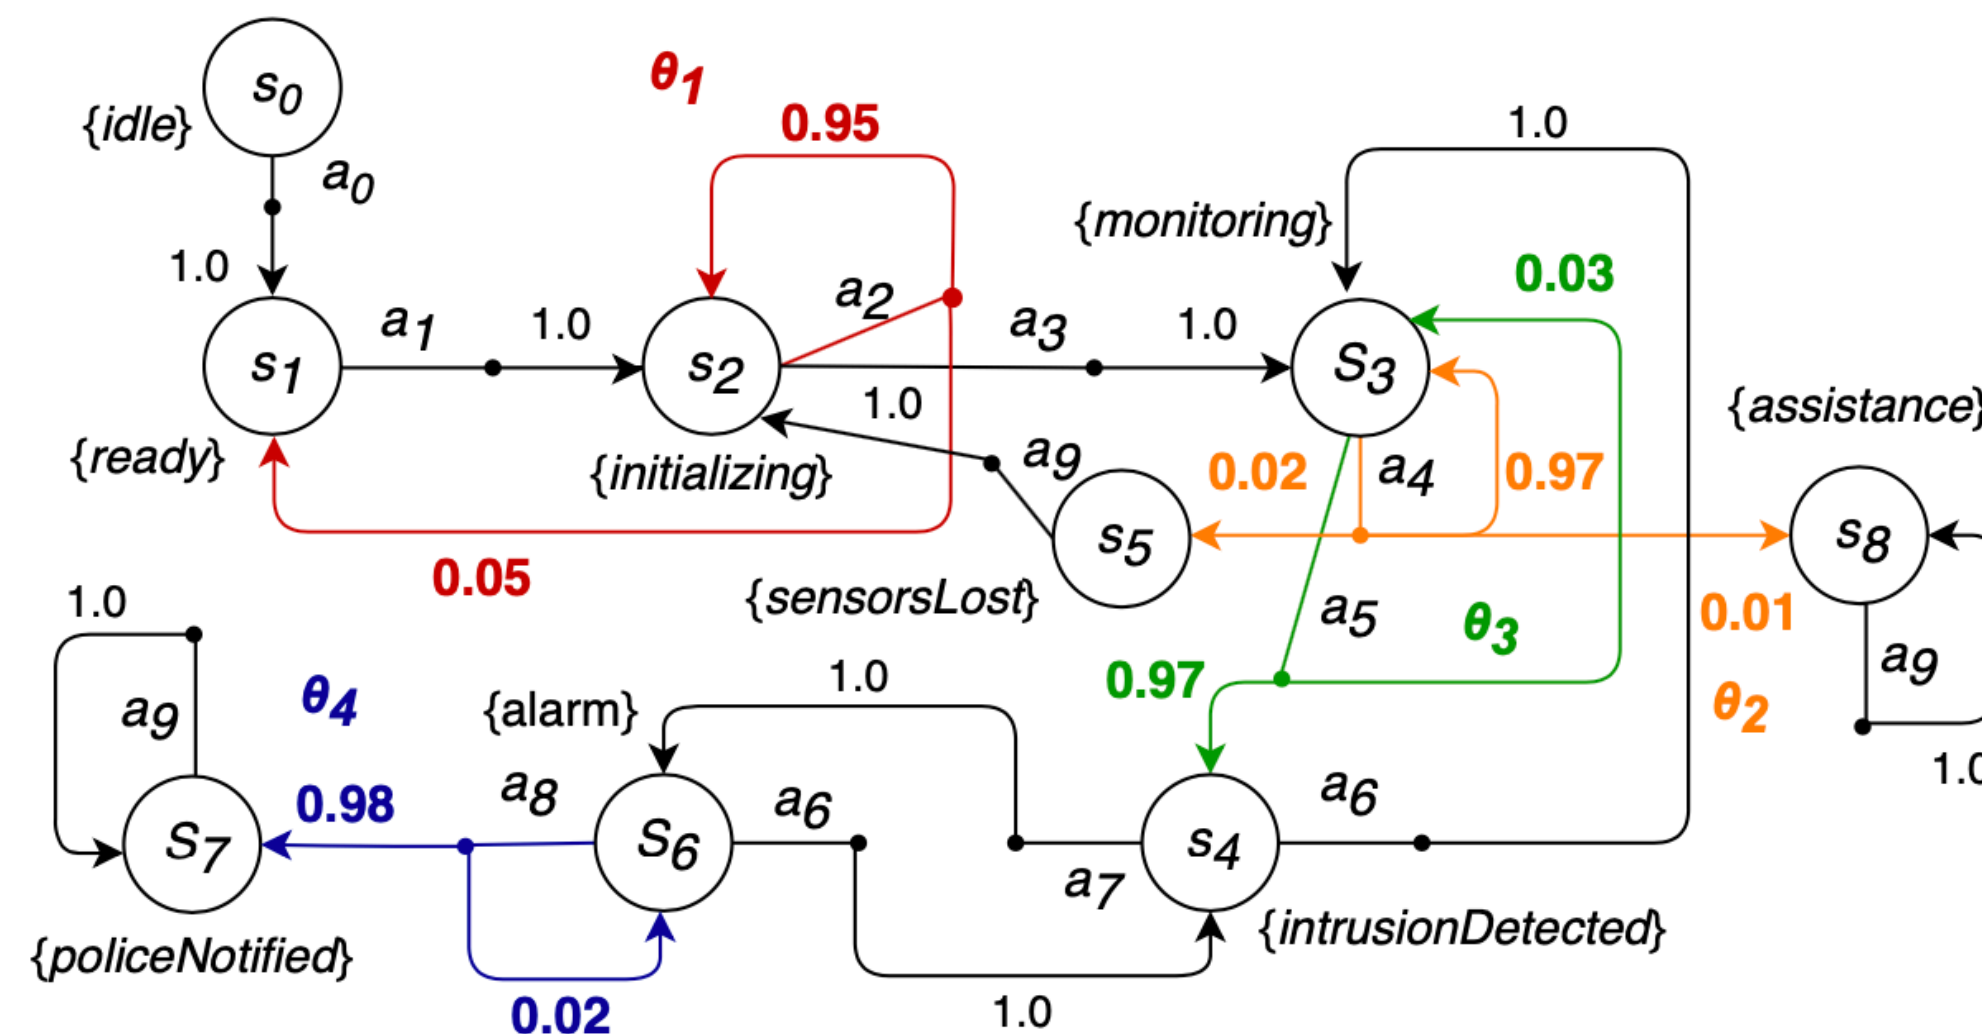
- Design-time models are imperfect and include assumptions
- sources of uncertainty introduce a mismatch between design-time assumptions and delivered software / production environment
 - Uncertain system properties (e.g., algorithmic/structural uncertainty, performance)
 - Uncertain environment properties (e.g., usage profiles, failure rate of 3rd-party components, latency)

The very idea

- **Objective**
 - Reduce the discrepancy between design-time assumptions (uncertain mathematical models) and real-world entities (runtime evidence)
- **How**
 - Mitigate sources of uncertainty by applying Inverse Uncertainty Quantification (IUQ)
 - Calibration of uncertain model parameters as soon as the system (or part of it) is available during integration/system testing by combining
 - Online Model-based Testing
 - Bayesian inference

Uncertain model parameters — SafeHome

- **Sources of uncertainty** (in CPSs) ¹
 - Application level — uncertain events/data originating from software running upon physical units (e.g., θ_3)
 - Infrastructure level — uncertain reliability of networking and/or cloud infrastructure (e.g., θ_4)
 - Integration level — uncertain outcomes from interacting physical units at either (e.g., θ_1, θ_2)



Examples

θ_2 : uncertain sensing capability from the monitoring state

θ_4 : uncertain failure rate of police notification from the alarm state

1. Man Zhang, et al., Uncertainty-wise cyber-physical system test modeling. Software & Systems Modeling (2017), 1–40

Uncertain regions

- Uncertain regions

- Uncertain transition probabilities θ_i grouped by $\langle \text{src-state}, \text{action} \rangle$
- Values in θ_i are uncertain parameters of a Categorical distribution
 - $\theta_i \sim \text{Cat}(p_1, \dots, p_k)$

region	state-action	affected level	target states	probability values
θ_1	$s_2 - a_2$	integration	s_2, s_1	0.95, 0.05
θ_2	$s_3 - a_4$	integration	s_3, s_4	0.03, 0.97
θ_3	$s_3 - a_5$	application	s_3, s_5, s_8	0.01, 0.97, 0.02
θ_4	$s_6 - a_8$	infrastructure	s_6, s_7	0.02, 0.98

Region

Uncertain parameters of a
Categorical distribution
—
hypothesis

- Intuition

- Mitigate the uncertainty over θ regions by observing the SUT
- Observation provides evidence to increase the confidence on transition probabilities

Bayesian inference

- Method used to update the probability for a hypothesis as more evidence becomes available
- **Formulation**¹
 - To learn θ (phenomenon of interest) we collect a sample $y = (y_1, \dots, y_n)$
 - $Posterior \propto likelihood \times Prior$
 - Prior $f(\theta)$ — hypothesis on θ
 - Likelihood $f(y \mid \theta)$ — compatibility of the evidence with the given hypothesis
 - Posterior $f(\theta \mid y)$ — best knowledge on the hypothesis given the evidence
- **In our context**
 - The natural conjugate Prior of the Categorical distribution is the Dirichlet distribution
 - $Prior_{\theta_i} \sim Dir(\alpha_1, \dots, \alpha_K)$
 - e.g., uninformative $Prior_{\theta_3} \sim Dir(0.5, 0.5, 0.5)$
 - e.g., informative $Prior_{\theta_3} \sim Dir(1.0, 97.0, 2.0)$
100 observations = 1 s_3 , 97 s_5 , 2 s_8

region	state-action	affected level	target states	probability values
θ_1	s_2-a_2	integration	s_2, s_1	0.95, 0.05
θ_2	s_3-a_4	integration	s_3, s_4	0.03, 0.97
θ_3	s_3-a_5	application	s_3, s_5, s_8	0.01, 0.97, 0.02
θ_4	s_6-a_8	infrastructure	s_6, s_7	0.02, 0.98

1. Robert, Christian. The Bayesian choice: from decision-theoretic foundations to computational implementation. Springer Science & Business Media, 2007

Bayesian inference (2)

• In our context

- Updating rule $\text{Prior}_{\theta_i} \sim \text{Dir}(\alpha_1, \dots, \alpha_K) \rightarrow \text{Post}_{\theta_i} \sim \text{Dir}(\alpha_1 + n_1, \dots, \alpha_K + n_K)$
 - e.g., $\text{Dir}(1.0, 97.0, 2.0) \rightarrow \text{Post}_{\theta_i} \sim \text{Dir}(1.0 + 35, 97.0 + 955, 2.0 + 10)$
1000 observations = 35 s_3 , 955 s_5 , 10 s_8

• Summarization

- Prior/Posterior knowledge can be summarized by using

- Mean — transition probability values $p_i = \alpha_i / \sum_{j=1}^k \alpha_j$

- HPD region — degree of confidence $C = \{p : f(\cdot) \geq 1 - \alpha\}, \alpha = 0.05$

region	state-action	affected level	target states	probability values
θ_1	s_2-a_2	integration	s_2, s_1	0.95, 0.05
θ_2	s_3-a_4	integration	s_3, s_4	0.03, 0.97
θ_3	s_3-a_5	application	s_3, s_5, s_8	0.01, 0.97, 0.02
θ_4	s_6-a_8	infrastructure	s_6, s_7	0.02, 0.98

Prior mean

inference

Posterior mean

0.033, 0.956, 0.011

Prior HPD region = { [0.009, 0.019], [0.936, 0.996], [0.001, 0.047] }

Prior HPD width = 0.136

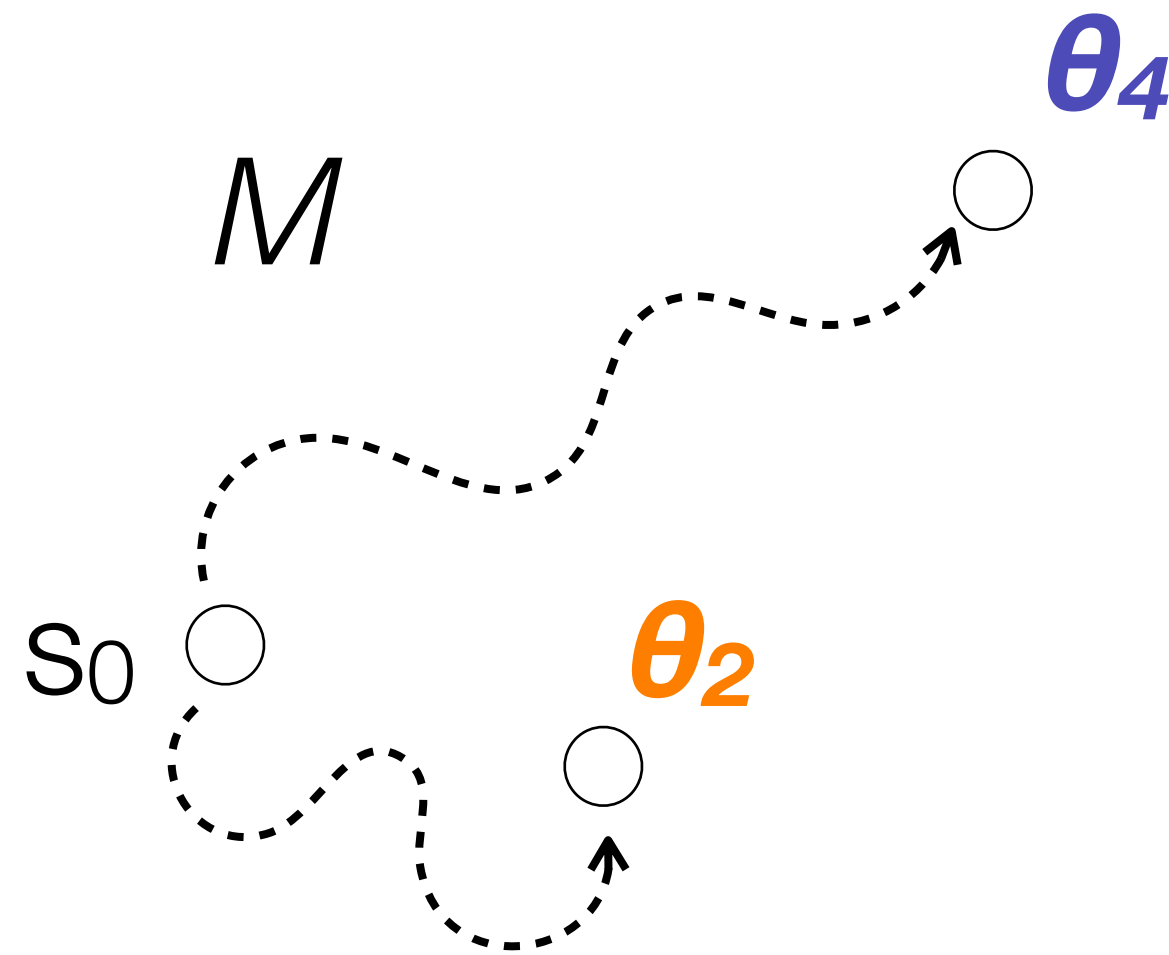
Posterior HPD region = { [0.009, 0.019], [0.936, 0.996], [0.001, 0.047] }

Posterior HPD width = 0.056

Lower value → higher confidence

Online MBT + Bayesian inference

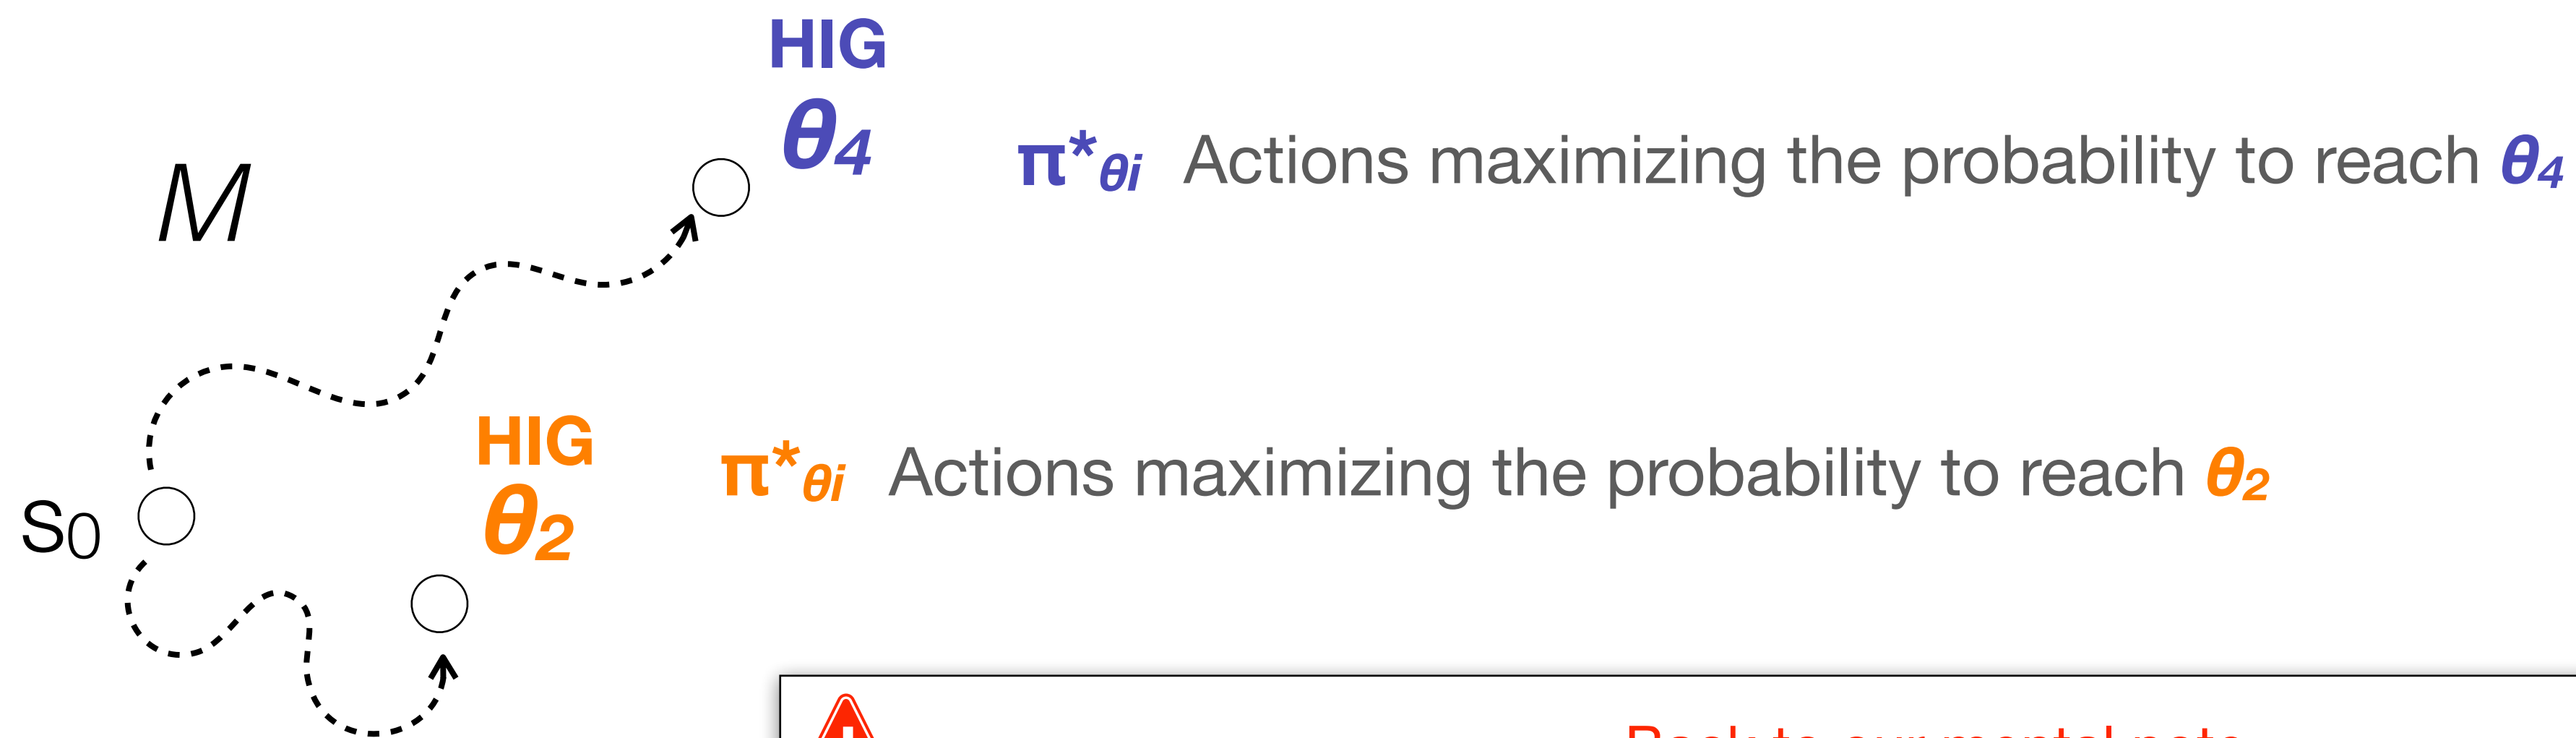
- **Objective** (reminder)
 - design-time assumptions \leftarrow gap reduction \rightarrow runtime evidence
- **How**
 - Perform a controlled exploration using online MBT to stress the uncertain components
 - Gather evidence and run bayesian inference to reduce the uncertainty



- **Uncertainty-aware MBT strategy**
 - Explore by maximize the probability of reaching θ regions
 - Reduces to an **optimization problem**:
 - Find out the actions a decision maker should take to maximize the exploration of θ regions

Uncertainty-aware strategy

- **Computation of the best policies**
 - For each θ_i
 - construct a reward structure that assigns **HIG** reward to θ_i transitions, **LOW** elsewhere
 - Compute the best policy $\pi^*_{\theta_i}$ (value iteration)
 - For each state, it selects the action that maximizes the probability to reach θ_i



Back to our mental note

We'll see how to leverage rewards in a "unconventional" way to drive testing.

Uncertainty-aware strategy (2)

- How to combine the best policies $\pi^*_{\theta_i}$?
 - Simple scenario \rightarrow there exists just a single θ region
 - Otherwise \rightarrow different exploration strategies may be constructed/adopted
 - Strategies represent decision makers (i.e., testers) that use a probabilistic function

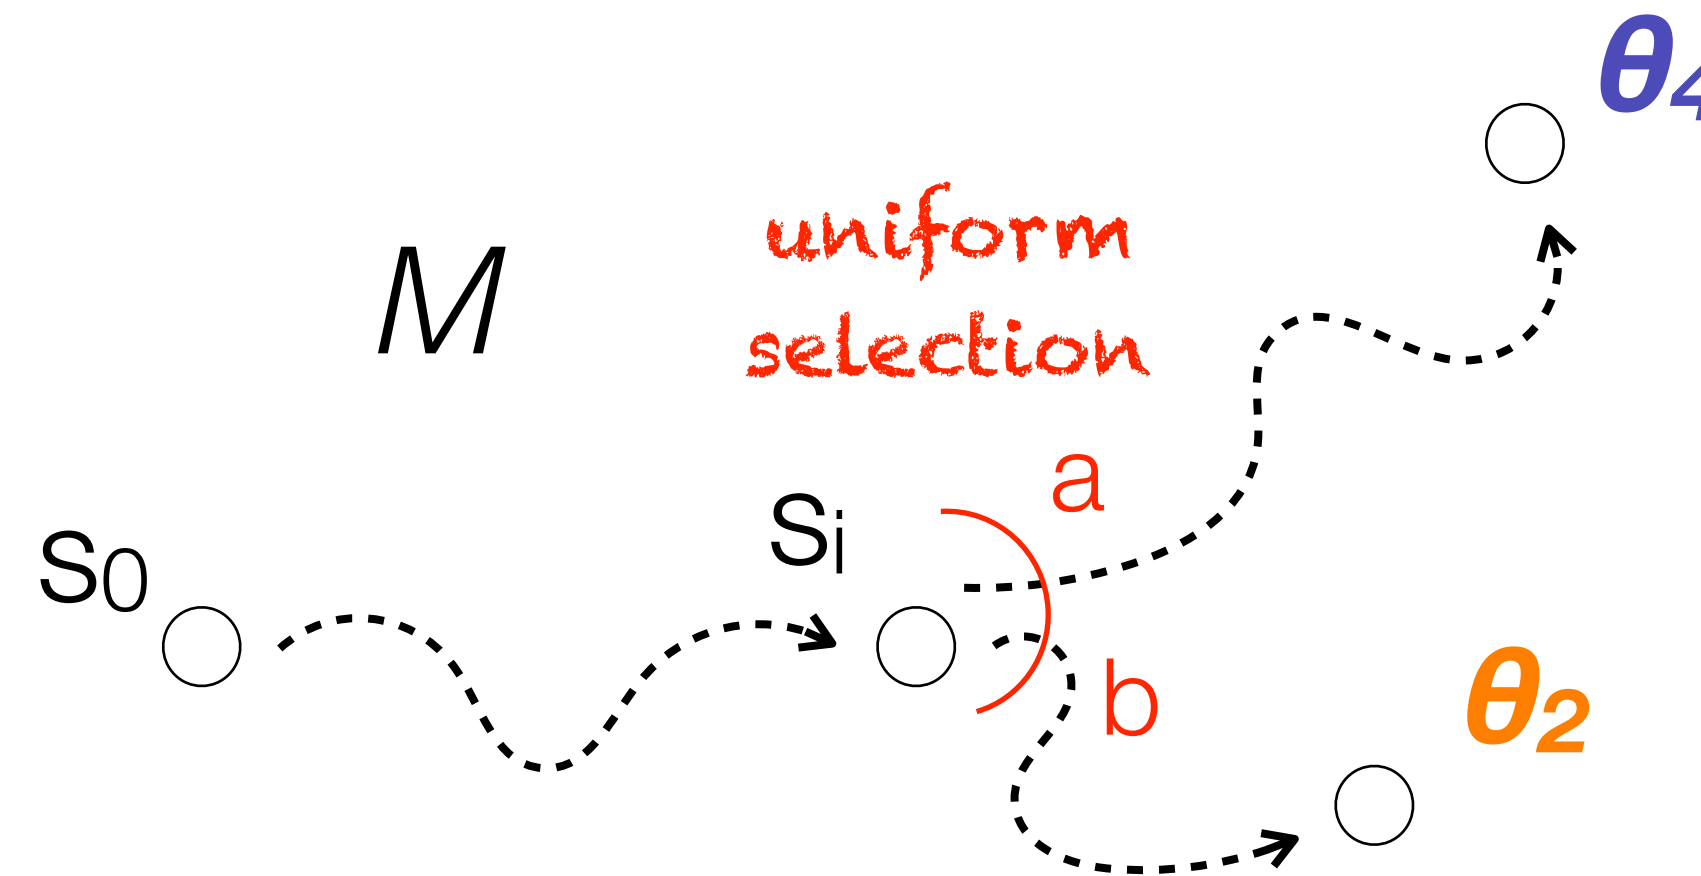
$$\mathcal{P}(s, a) = \begin{cases} 0 & \omega(s, a) = 0 \\ \omega(s, a) / \sum_{a' \in A(s)} \omega(s, a') & \text{otherwise} \end{cases}$$

- The ω weight selectively increase/decrease the probability of choosing a specific action

Uncertainty-aware strategy (3)

- **Flat strategy**
 - Actions selected by different policies $\pi^*_{\theta_i}$ have equal probability
 - Uniform random sampling of the available policies

$$\omega^{RT}(s, a) = \begin{cases} 1 & \exists i : \pi_i^*(s) = a \\ 0 & \text{otherwise} \end{cases}$$

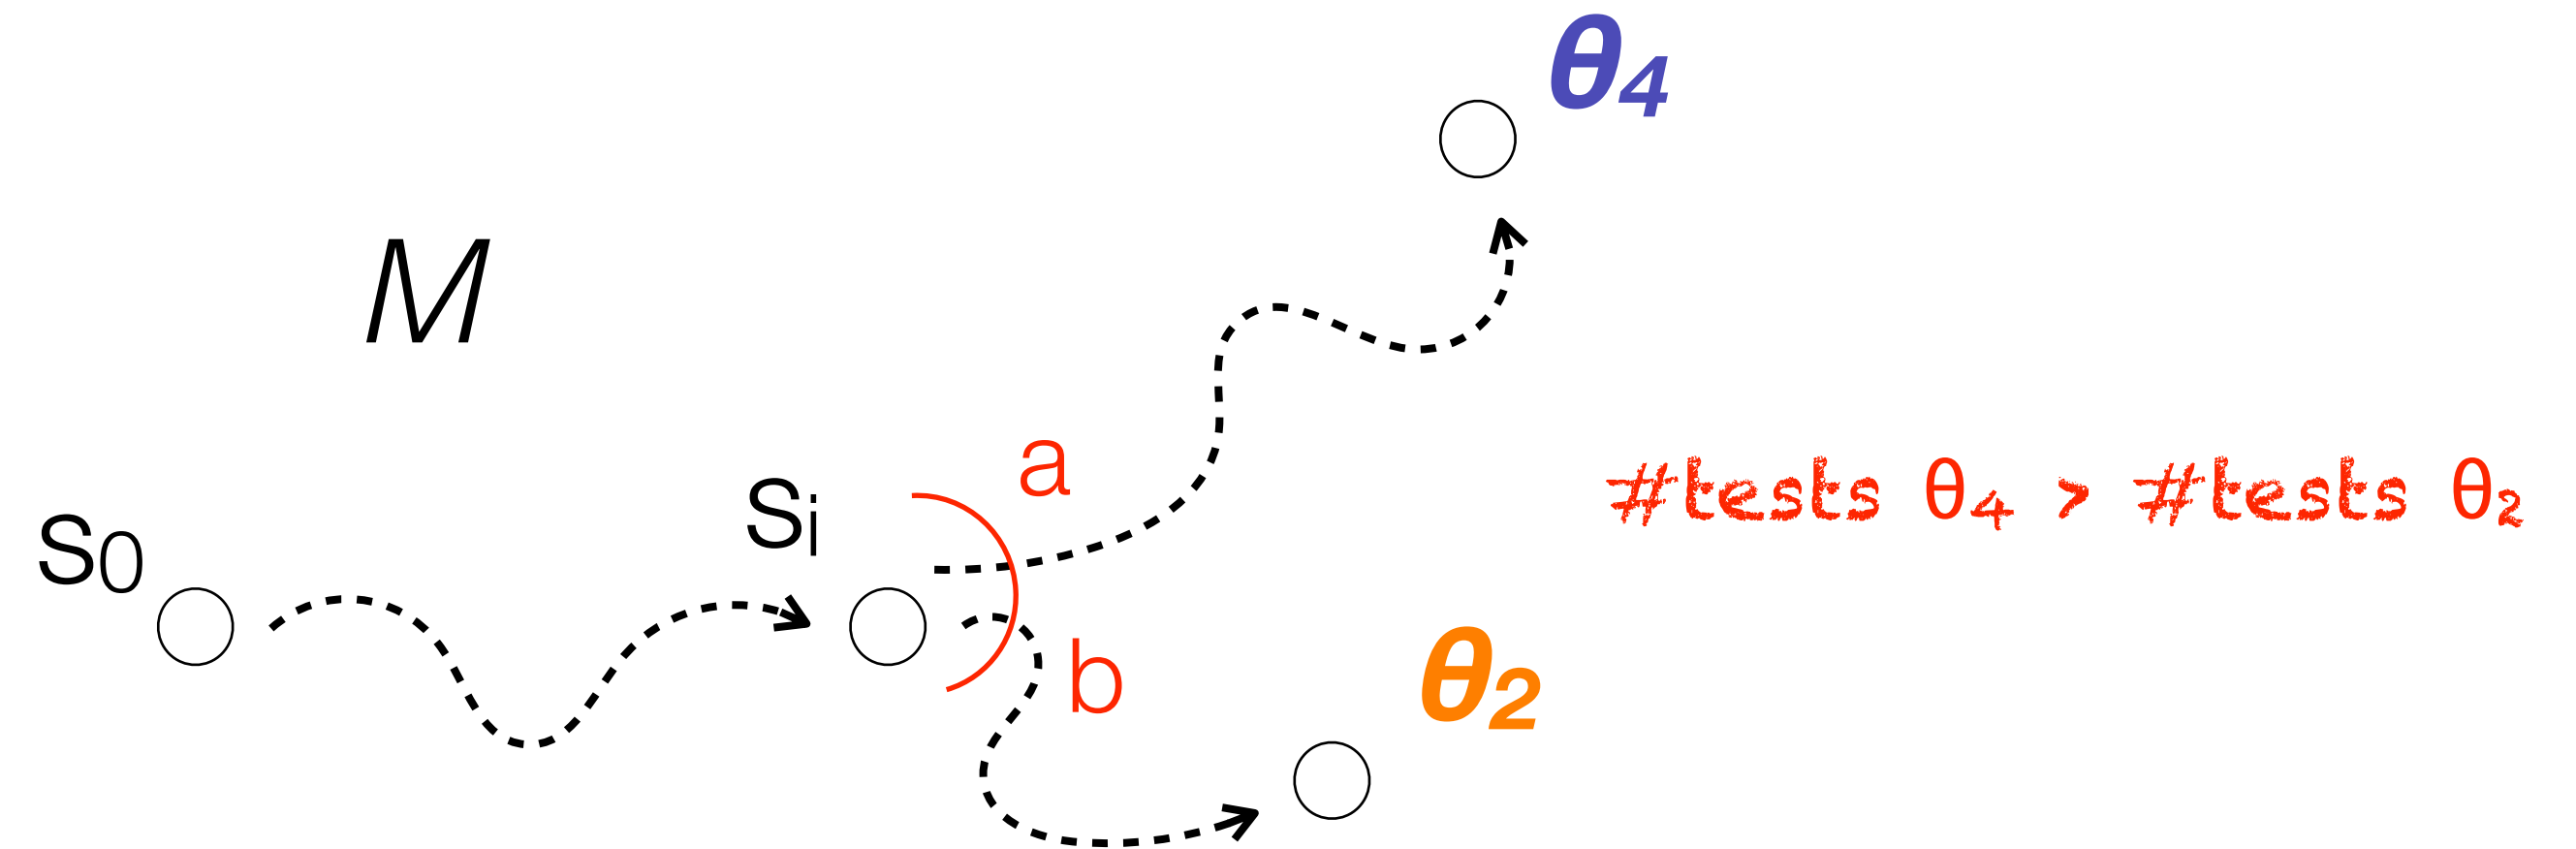


Uncertainty-aware strategy (4)

- History-based strategy

- Tries to keep balanced the number of times θ regions are tested
- We leverage decrementing weights inversely proportional to #selections of state-action pairs

$$\omega^{HT}(s, a) = \begin{cases} 1/\#(s, a) & \exists i : \pi_i^*(s) = a \\ 0 & \text{otherwise} \end{cases}$$

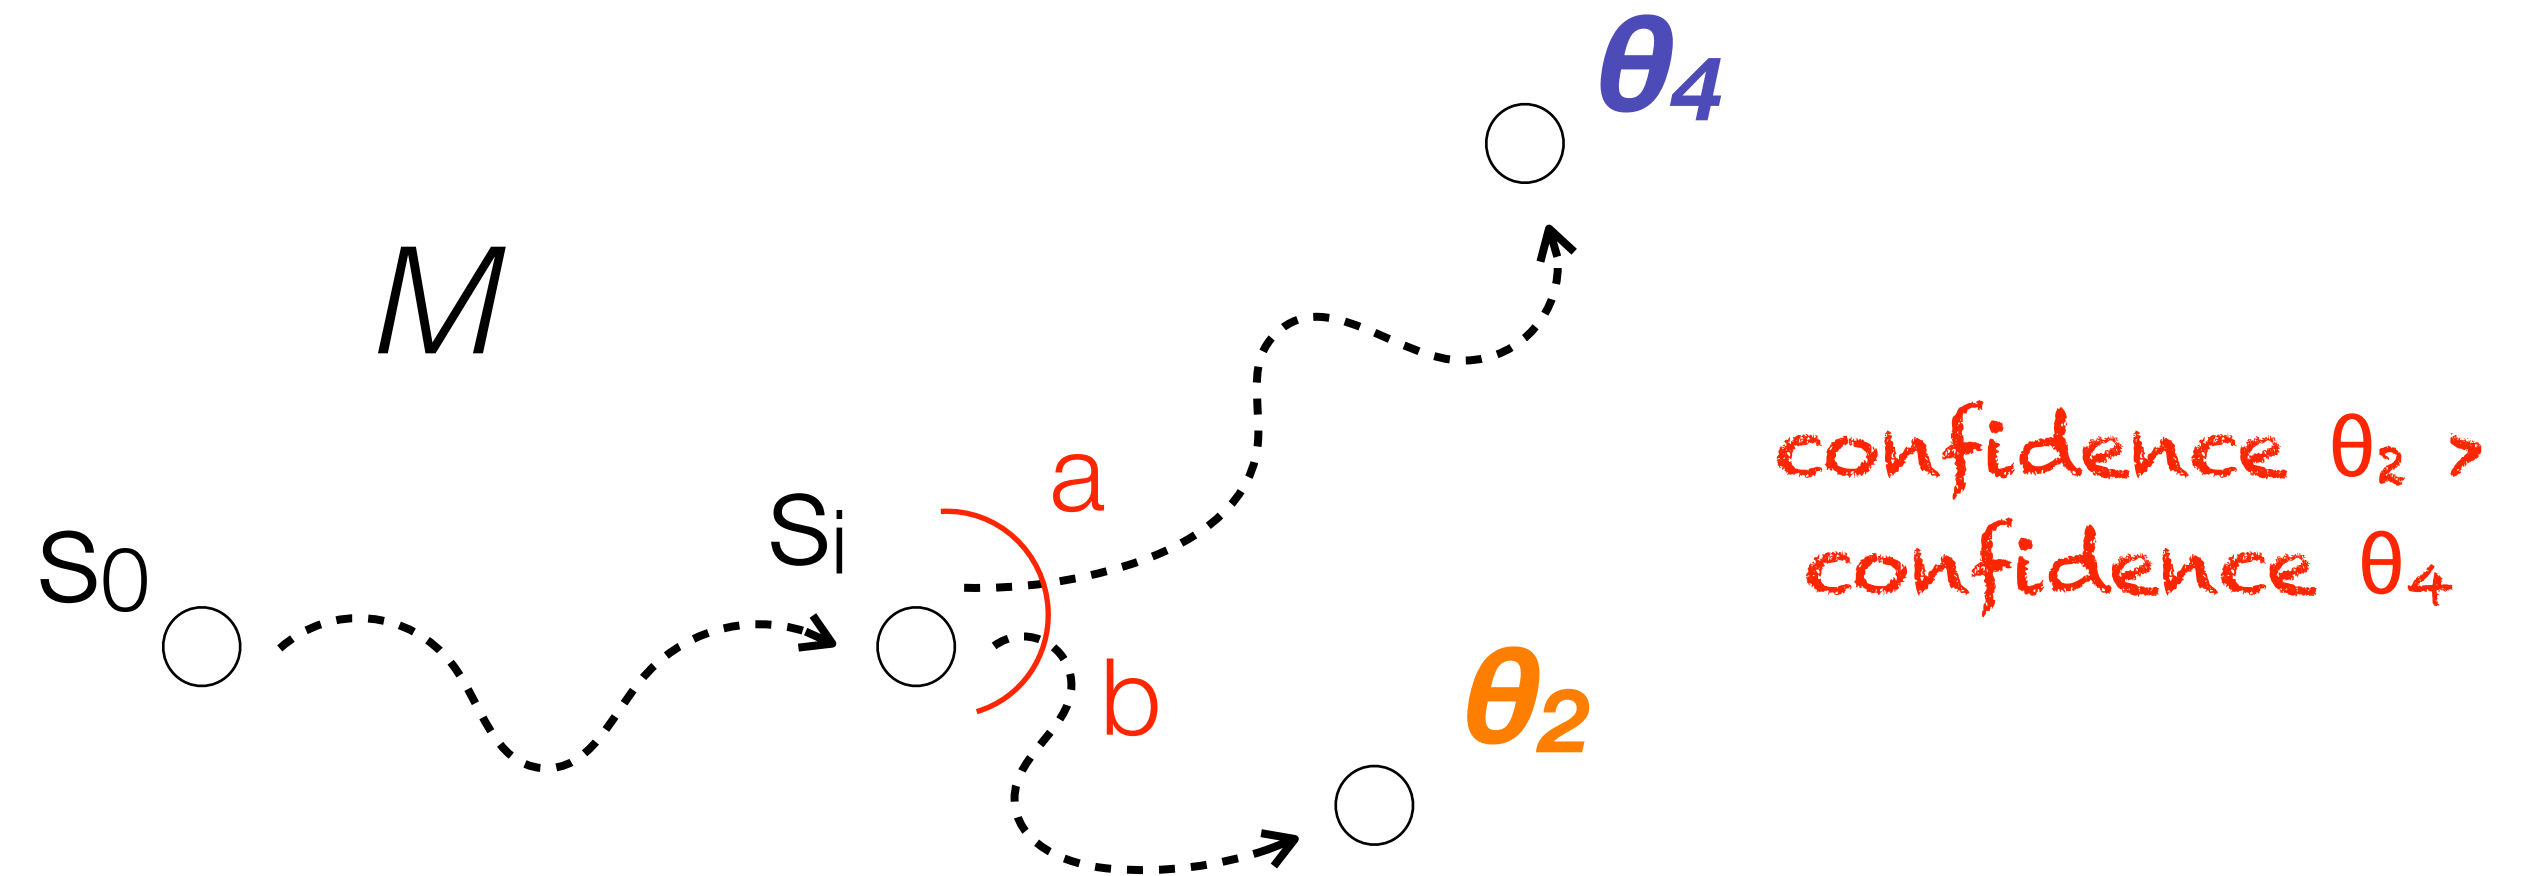


the higher the #selections of a,
the lower the likelihood of selecting it again

Uncertainty-aware strategy (5)

- **Distance strategy**
 - Tries to deliver balanced degree of confidence on θ regions
 - The weight is proportional to the HPD width of θ regions

$$\omega^{DT}(s, a) = \begin{cases} |\text{HPD}_{\theta_i}| & \exists i : \pi_i^*(s) = a \\ 0 & \text{otherwise} \end{cases}$$



the larger the HPD width of the target θ ,
the higher the likelihood of selecting it

Termination condition

- Limit on the effort
 - Traditional termination condition based on #tests limit
- Bayes factor
 - Tries to recognize when the inference process converges

$$\mathcal{F} = \frac{f(y|\theta)}{f(y|\theta')}$$

Likelihood that data y are produced
under different assumptions θ and θ'

$$\mathcal{F} \in [10^0, 10^{1/2}]$$

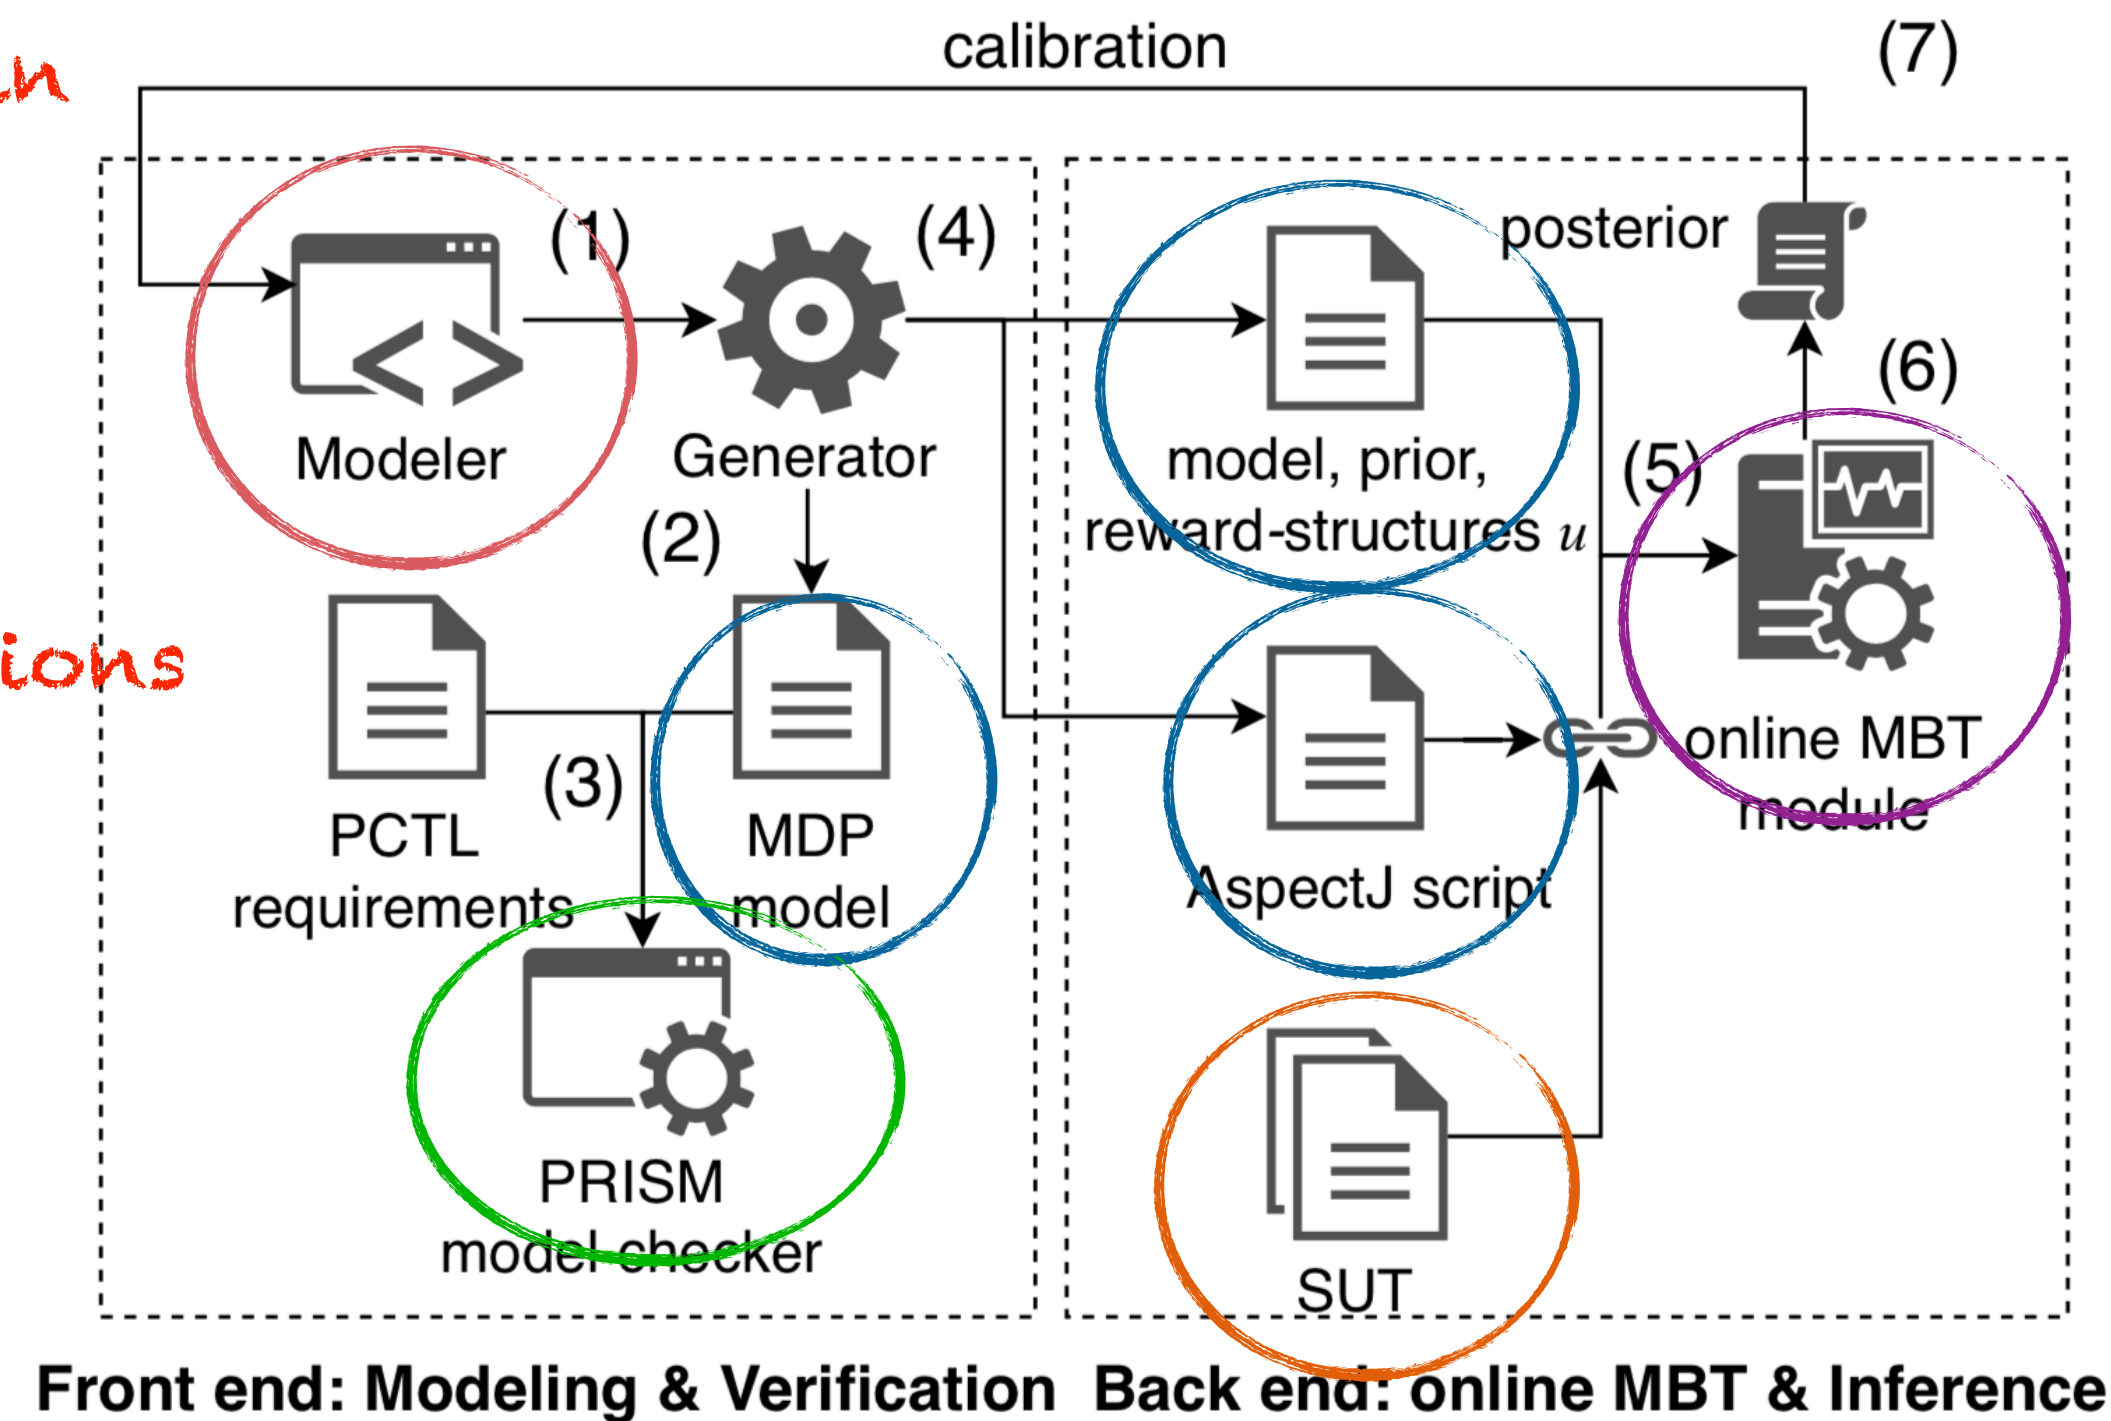
Difference between assumptions θ and θ'
is not substantial

Current toolchain implementation

MDP definition + uncertain
transition probabilities θ

binding to the SUT:
actions \rightarrow inputs
acs \rightarrow routine postconditions

automatic generation



online MBT
and inference/
calibration

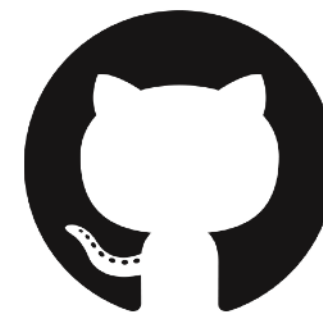
model checking of
PCTL requirements

Java program

Current toolchain implementation (2)

MBT module

<https://github.com/SELab-unimi/mbt-module>



● Java 99.3%
● Other 0.7%

- **Current stage**
 - Uncertainty-aware strategies have been implemented
 - Systematic evaluation of their cost-effectiveness —> currently inedited

Current toolchain implementation (3)

- **Evaluation summary**

- We assessed statistical difference (Mann-Whitney U test ¹)
- We evaluated practical value (Vargha & Delaney's \hat{A}_{12} measure ¹)
 - In our context \rightarrow assuming same effort (i.e., #tests), the probability that target strategy yields smaller HPD width values than flat one (i.e., baseline)

Table 3: Vargha and Delaney's \hat{A}_{12} measure

	%uncertainty			#actions		
balanced	20	50	80	5	10	20
hist	1.000	0.716	0.531	0.617	0.704	0.926
dist	1.000	0.790	0.679	0.741	0.802	0.951
unbalanced	20	50	80	5	10	20
hist	0.963	0.716	0.556	0.531	0.642	0.901
dist	0.988	0.951	0.691	0.741	0.741	0.975

1. Andrea Arcuri and Lionel Briand, A practical guide for using statistical tests to assess randomized algorithms in software engineering, ICSE'11, New York, NY, USA

Summary


- We discussed MBT for probabilistic systems and the problem of testing with uncertain model components
- Depending on the Prior knowledge (hypothesis) and information that can be gathered during testing, we derived different uncertainty-aware exploration strategies and evaluated their cost-effectiveness
 - Flat —> uniform selection
 - History —> balanced exploration
 - Distance —> balanced delivered confidence
- **Next**
 - Hands on session with the MBT module
 - Design/develop an additional exploration strategy

References

- Camilli M., Gargantini A., Scandurra P., Bellettini C. (2017) Towards Inverse Uncertainty Quantification in Software Development (Short Paper). In: Cimatti A., Sirjani M. (eds) Software Engineering and Formal Methods. SEFM 2017. LNCS, vol 10469. Springer, Cham.
- M. Camilli, C. Bellettini, A. Gargantini and P. Scandurra, (2018) Online Model-Based Testing under Uncertainty, IEEE 29th International Symposium on Software Reliability Engineering (ISSRE), Memphis, TN, 2018, pp. 36-46.
- M. Camilli, A. Gargantini, R. Madaudo and P. Scandurra, (2019) HYPpOTesT: Hypothesis Testing Toolkit for Uncertain Service-based Web Applications. In proceeding of the 15th International conference on integrated Formal Methods. iFM2019. LNCS, vol 11918. Springer, Cham.
- M. Camilli, A. Gargantini, and P. Scandurra, (2020) Model-based Hypothesis Testing of Uncertain Software Systems, Software Testing Verification and Reliability, John Wiley & Sons, Ltd. <https://doi.org/10.1002/stvr.1730>, To appear.

License of these slides

© 2019-2020 Matteo Camilli



Except where otherwise noted, this work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

To view a copy of this license, visit

<https://creativecommons.org/licenses/by-nc-nd/4.0/>